

Functional Infrastructures

Toni Batchelli, @disclosure

PalletOps

Clojure/West 2013



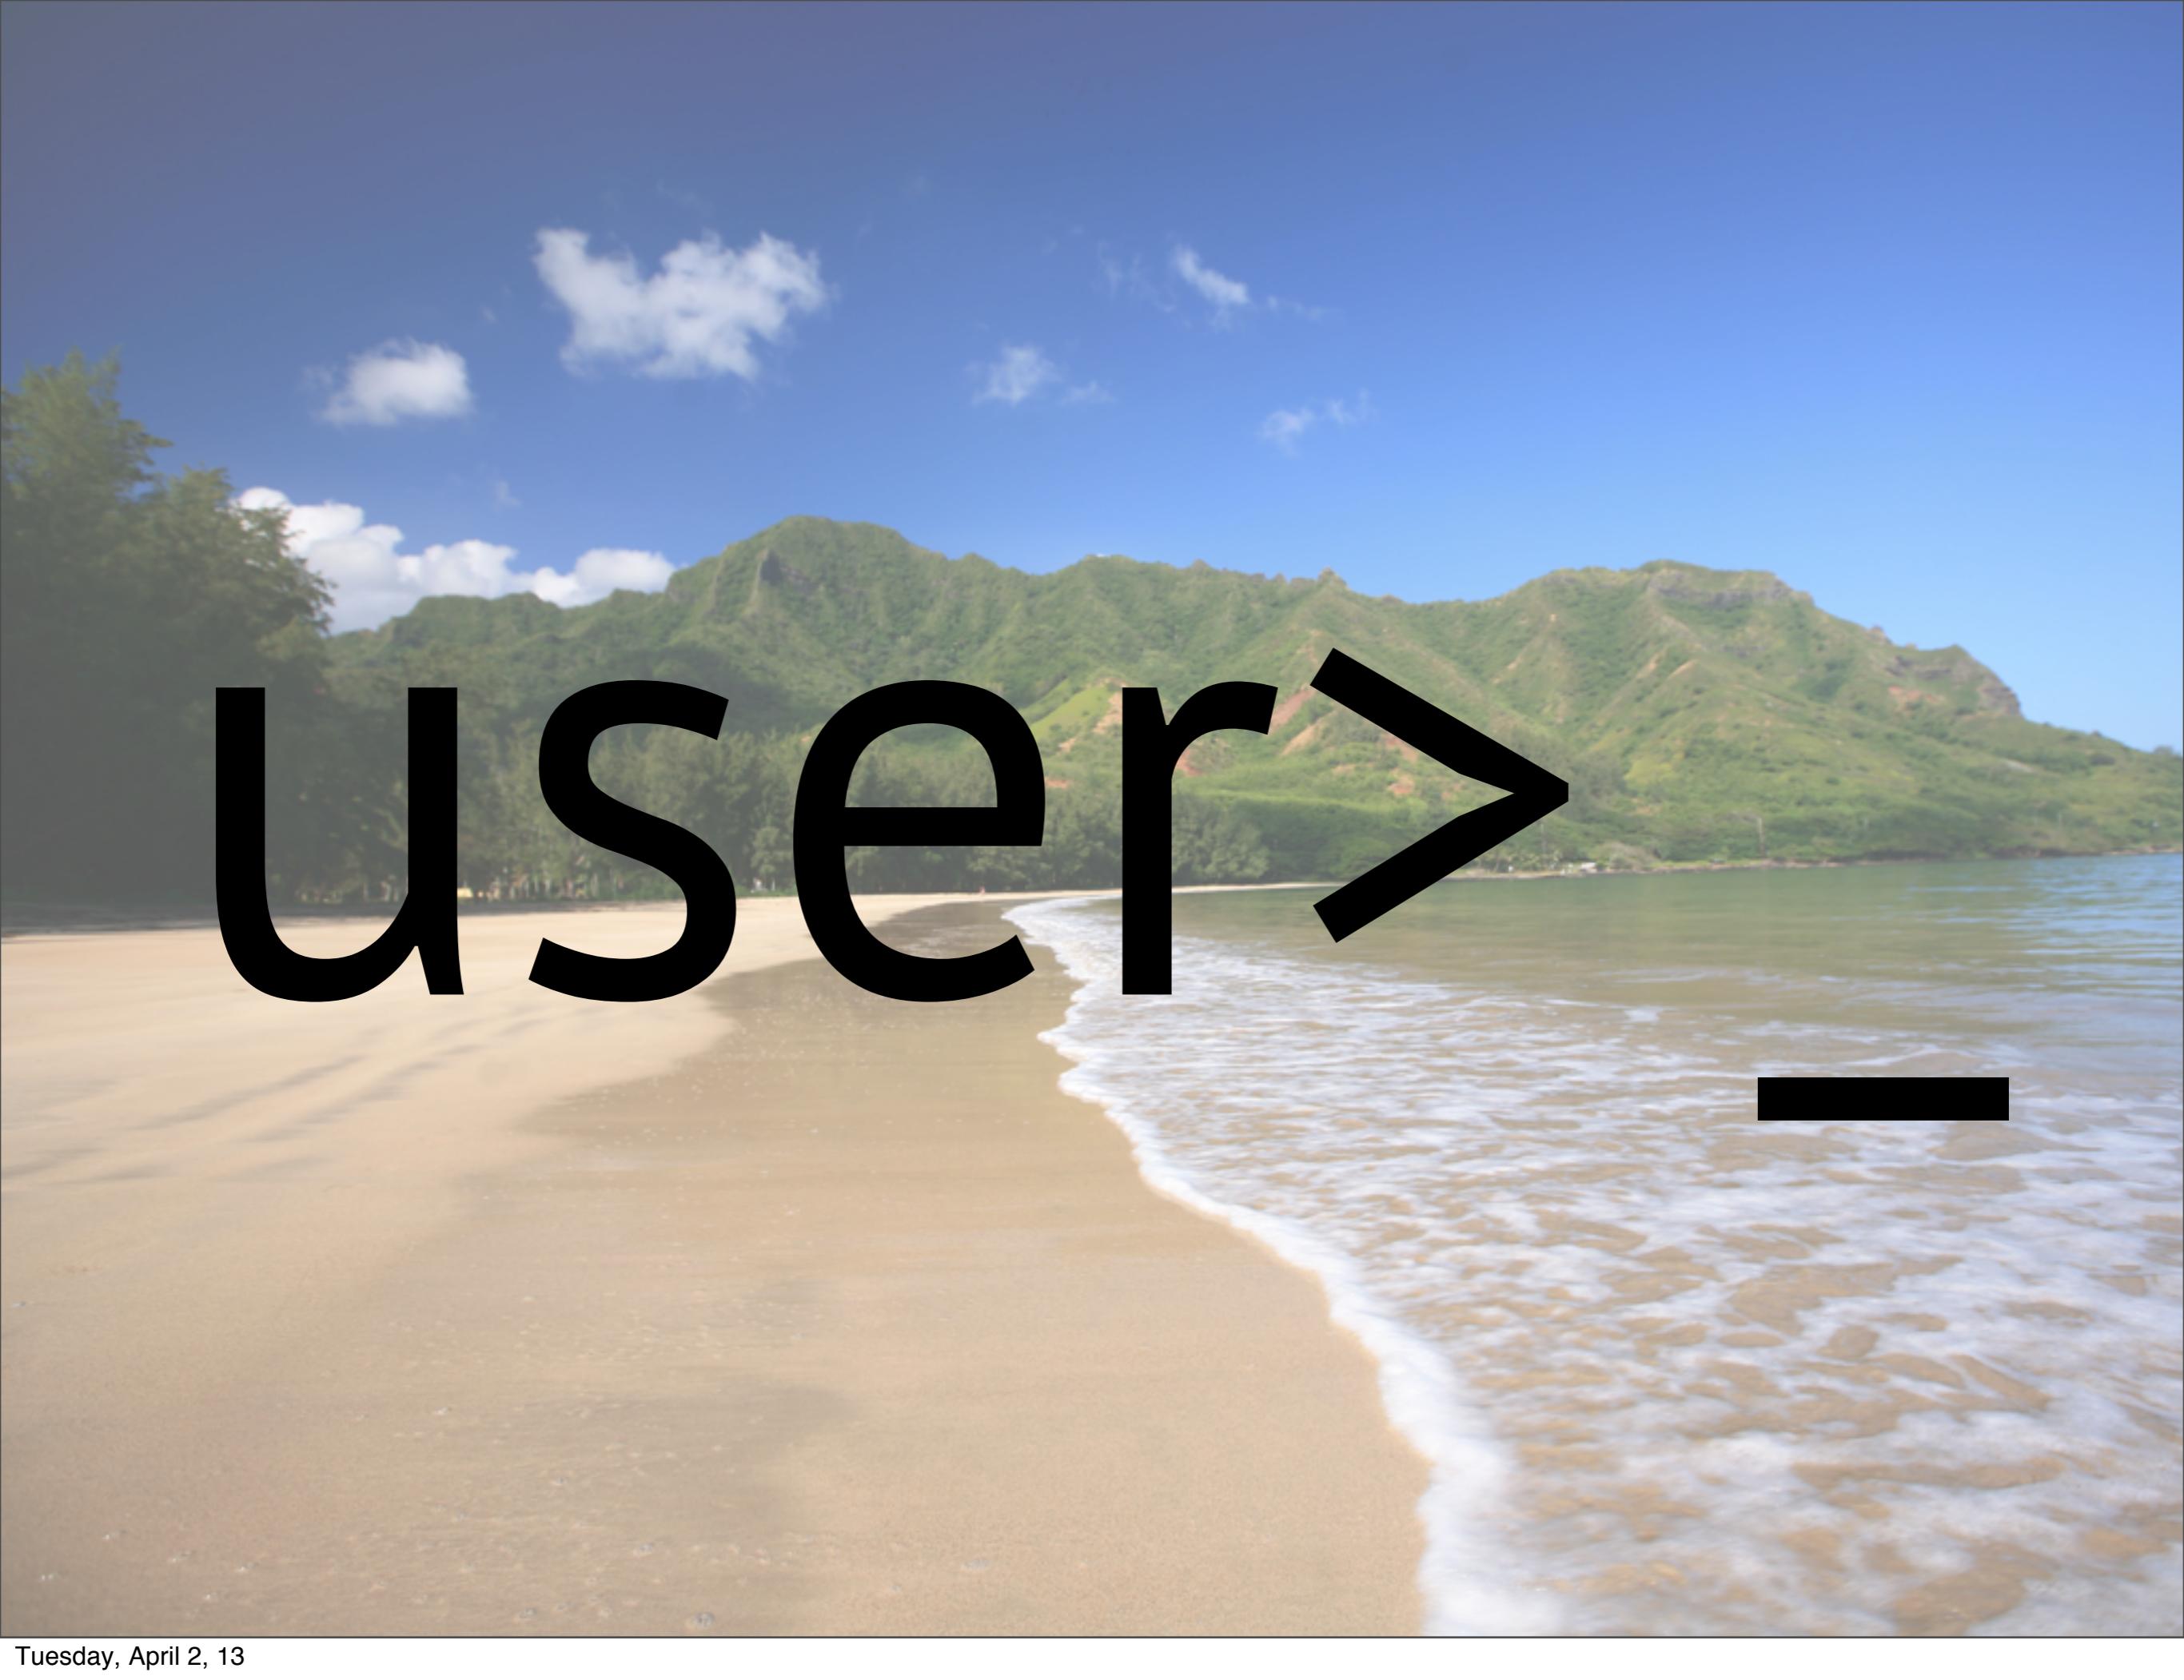
or...

**It's all fn until you hit
production**



user>



A photograph of a tropical beach with golden sand and gentle waves. In the background, there are two large, lush green mountains under a blue sky with scattered white clouds.

user>

[REDACTED]

b = **f(a)**

bash \$



bash\$



$b \approx f(a) ?$

b=rand(a)

Infrastructure Automation



desired
system

$-f(\text{current system})$

Bottom Up: Ops' Perspective



One box

- **Concrete: known hardware, OS family and OS version, package version, tooling**
- **Stateful (erk!)**



- **Scripts**
- **Save/Clone Image**

Many boxes

- Some configuration is shared
- Certain uniformity is desired



- Scripts+Templates
- Configuration Data Base
- Transport
- Declarative actions
- Image Save/Clone

Many boxes

- Some configuration is shared
- Certain uniformity is desired



- Scripts+Templates
- Configuration Data Base
- Transport
- Declarative actions
- Image Save/Clone

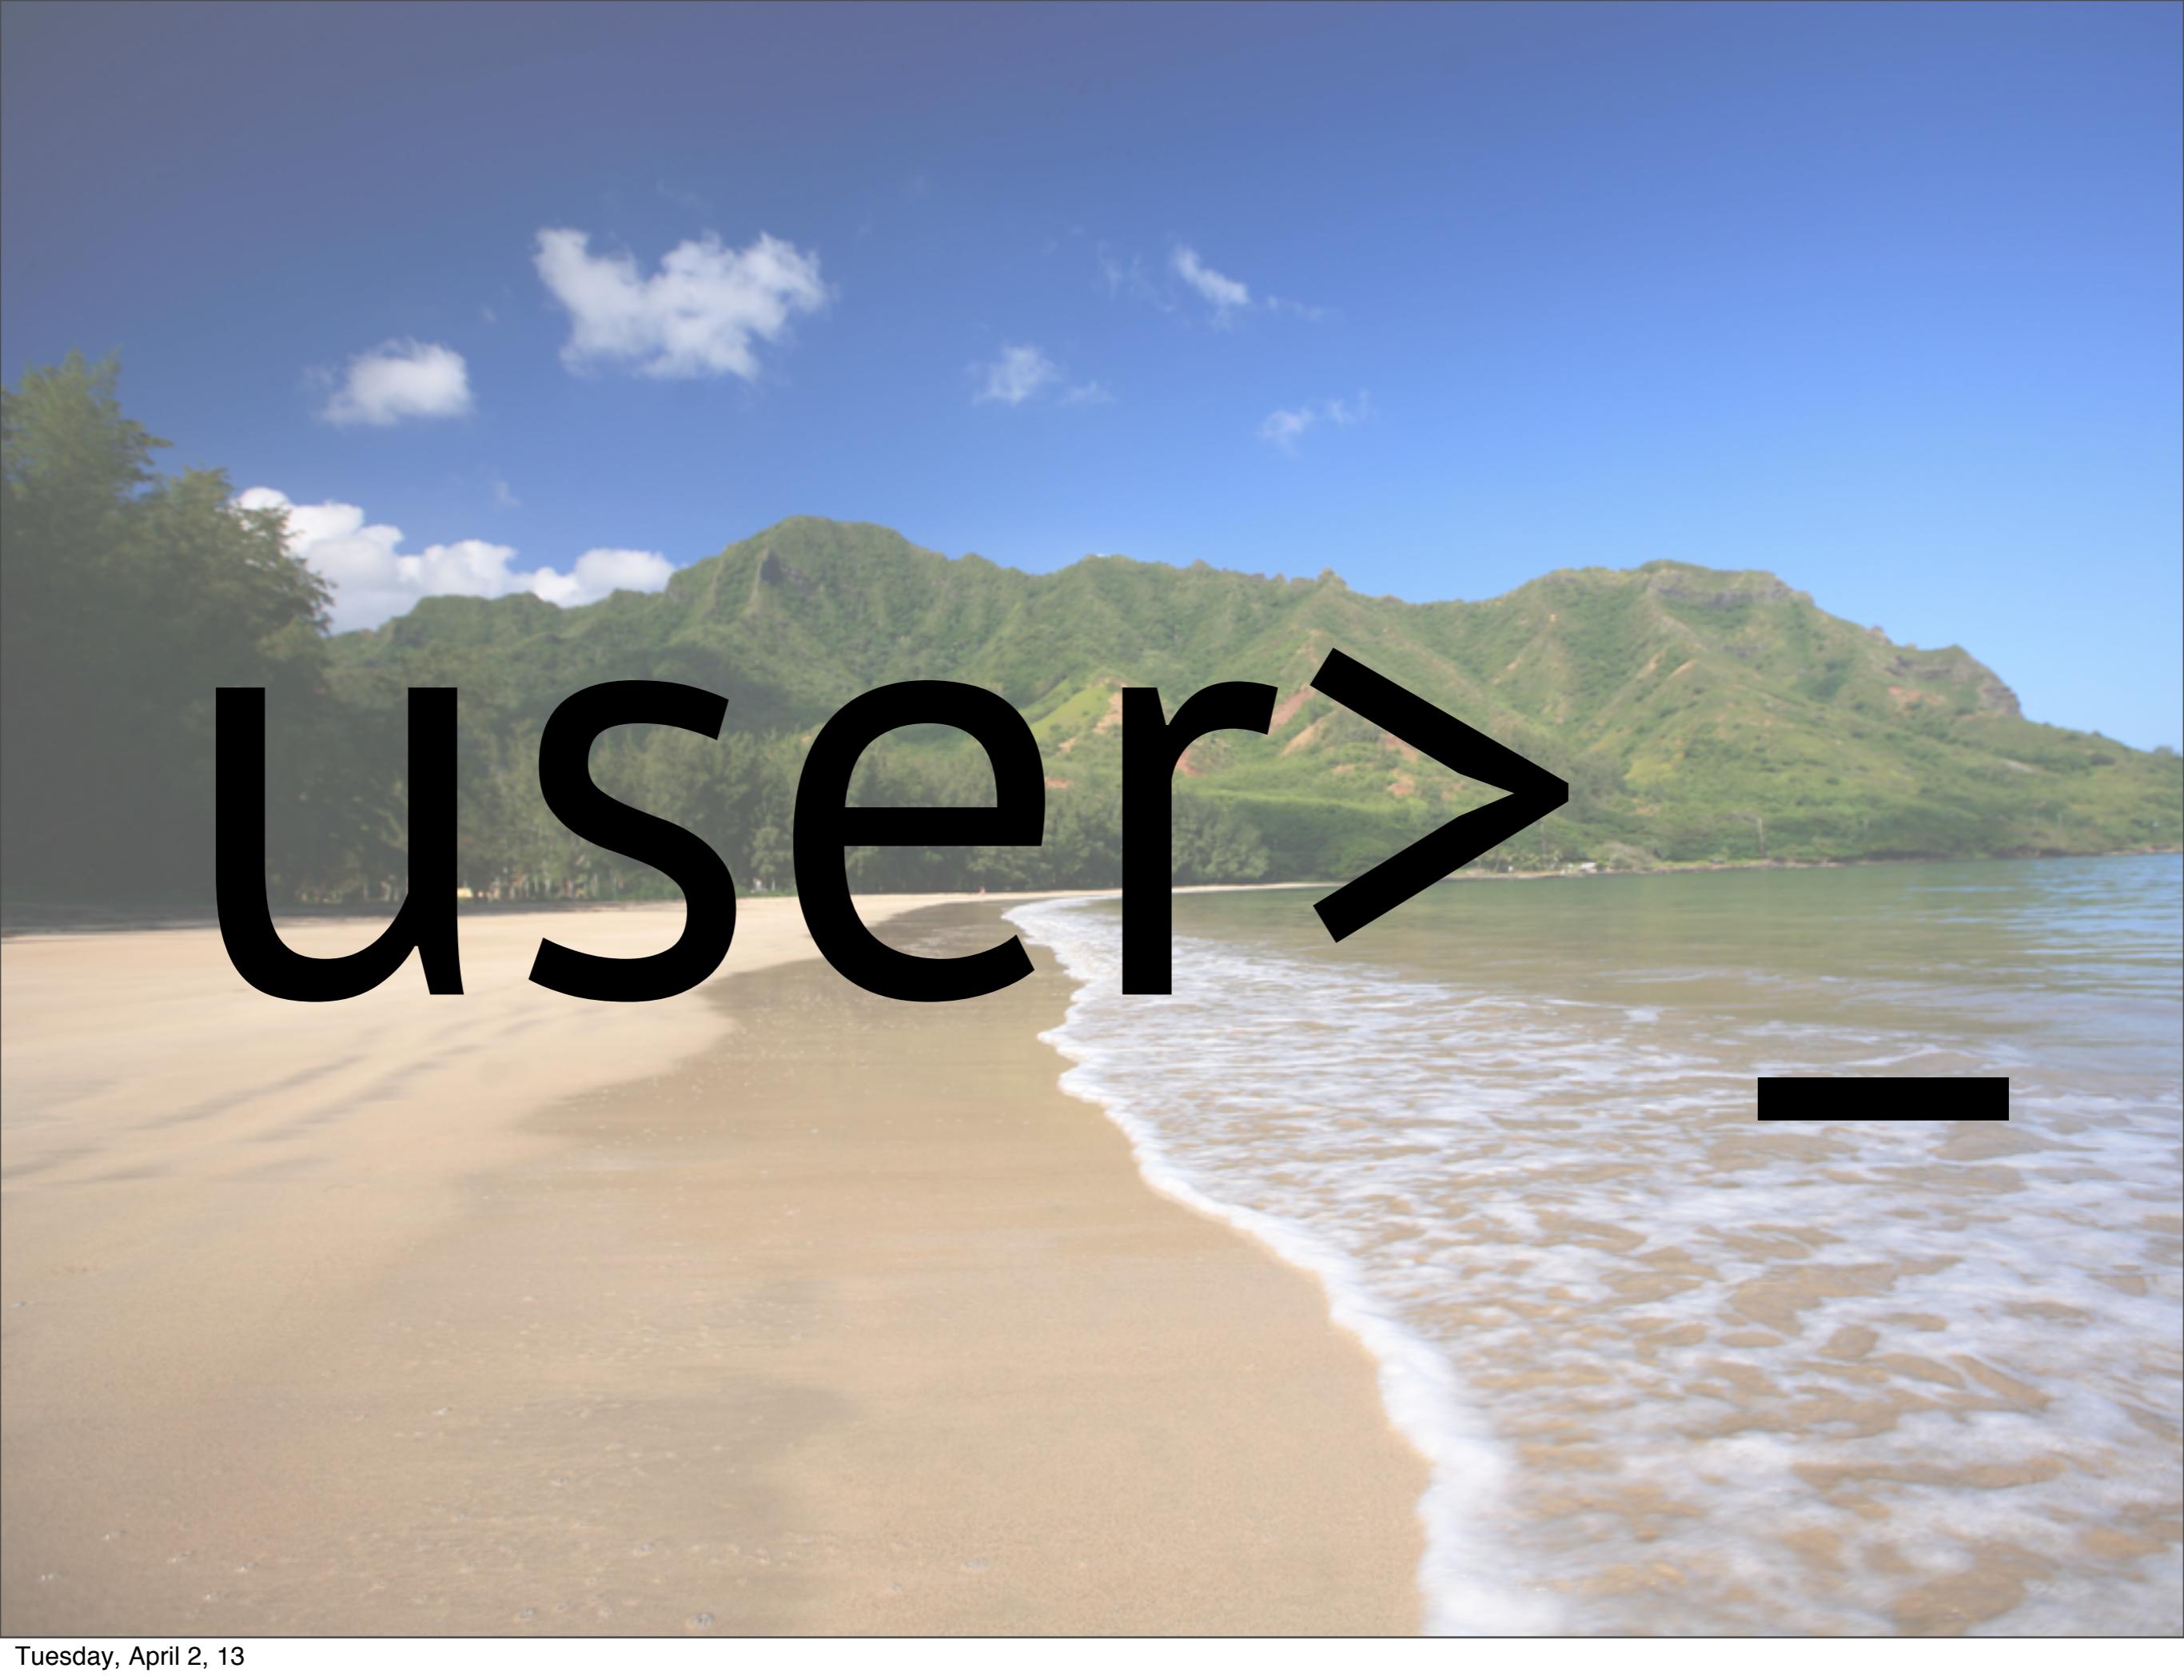


DevOps



Top Down: Dev's Perspective



A photograph of a tropical beach with golden sand and gentle waves. In the background, there are two large, lush green mountains under a blue sky with scattered white clouds.

user>

[REDACTED]

Dev's view

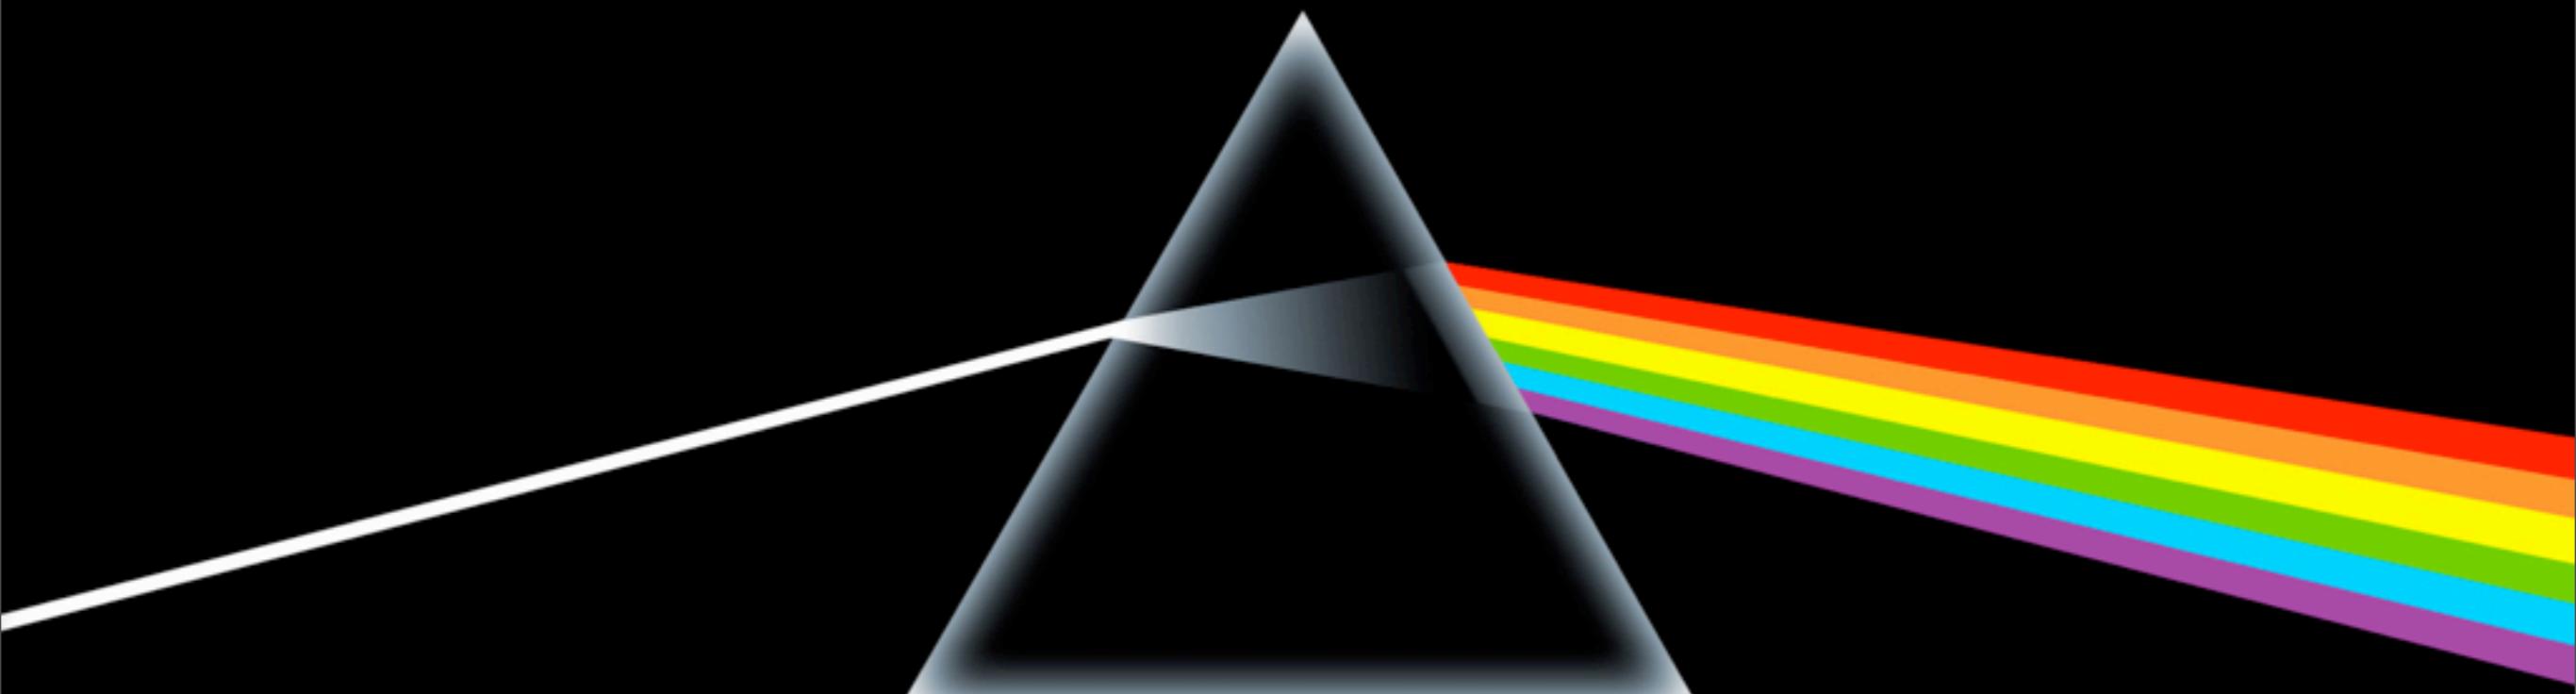
- OS & Hardware don't matter
- Size doesn't matter
- Build environments
 - Dev, Test, Stage & Production
- Same code everywhere
- Development branches



Dev's OTHER view

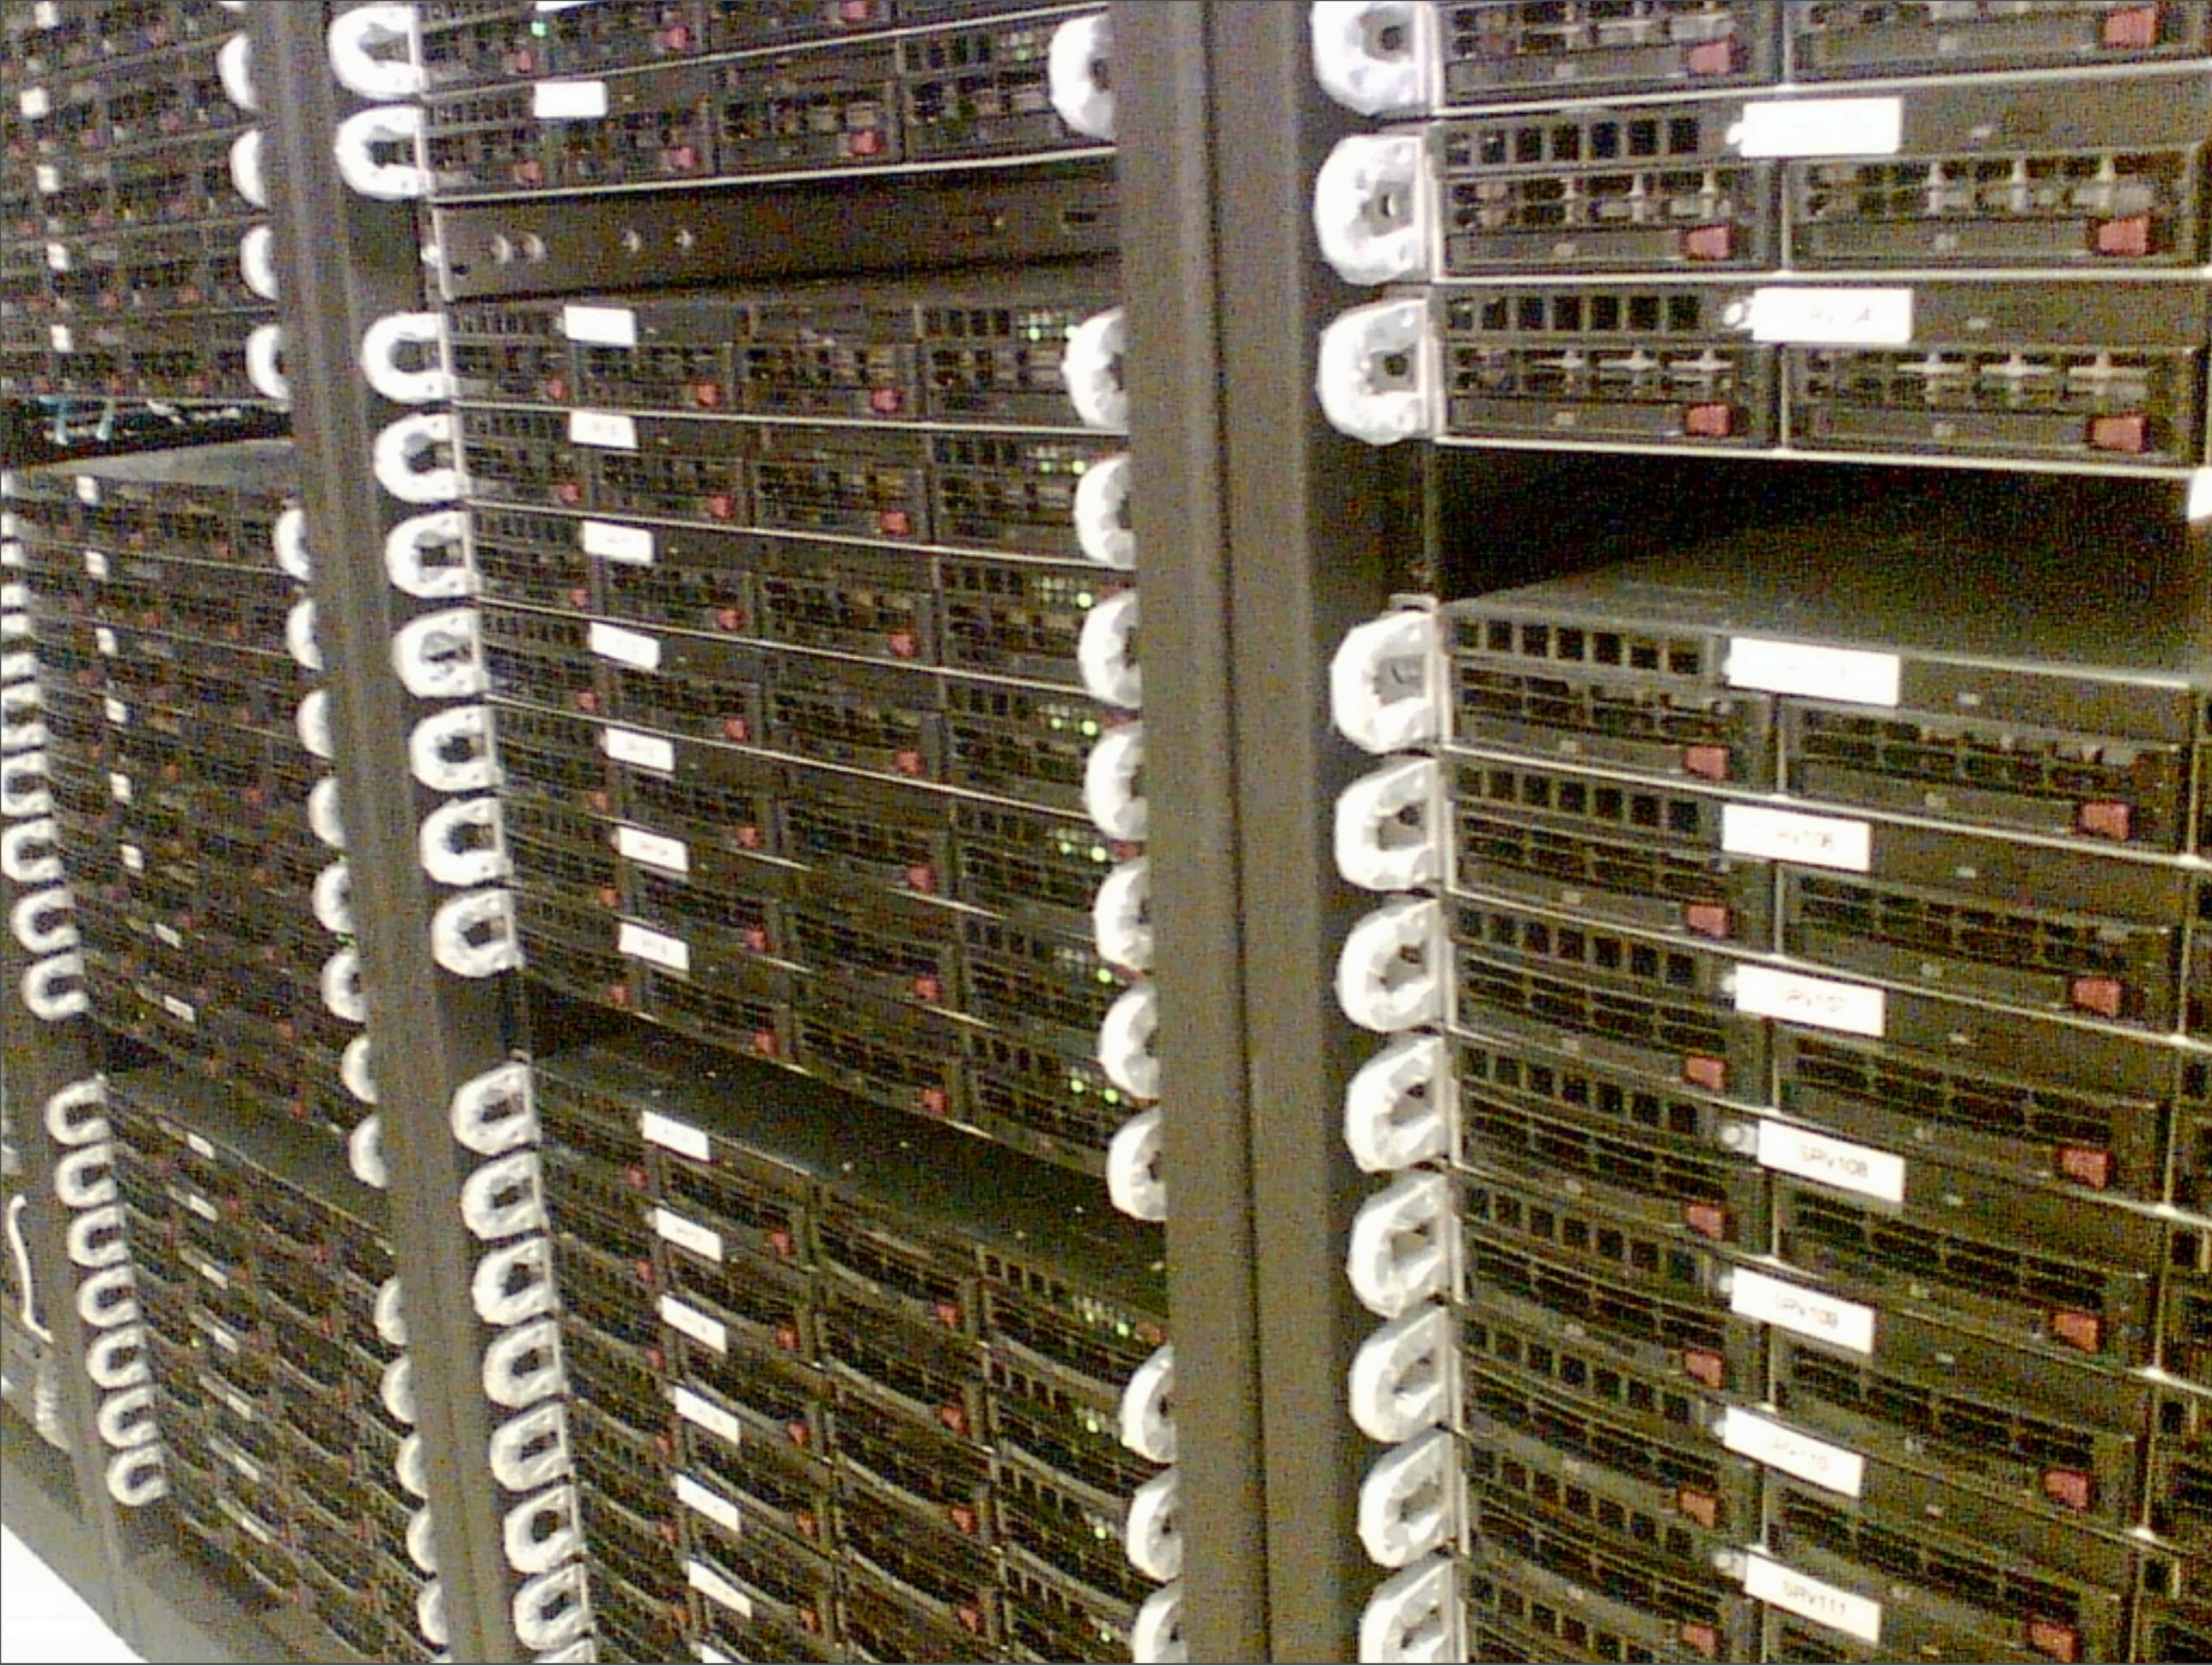
- Test on ubuntu 10.04 and Centos 6.3
- Prepare changes for new version of HBase
- Test recovery with hot stand-by DB
- Benchmark horizontal scalability
- Swap MySQL for MongoDB
- How do firewall rules affect our distributed system?



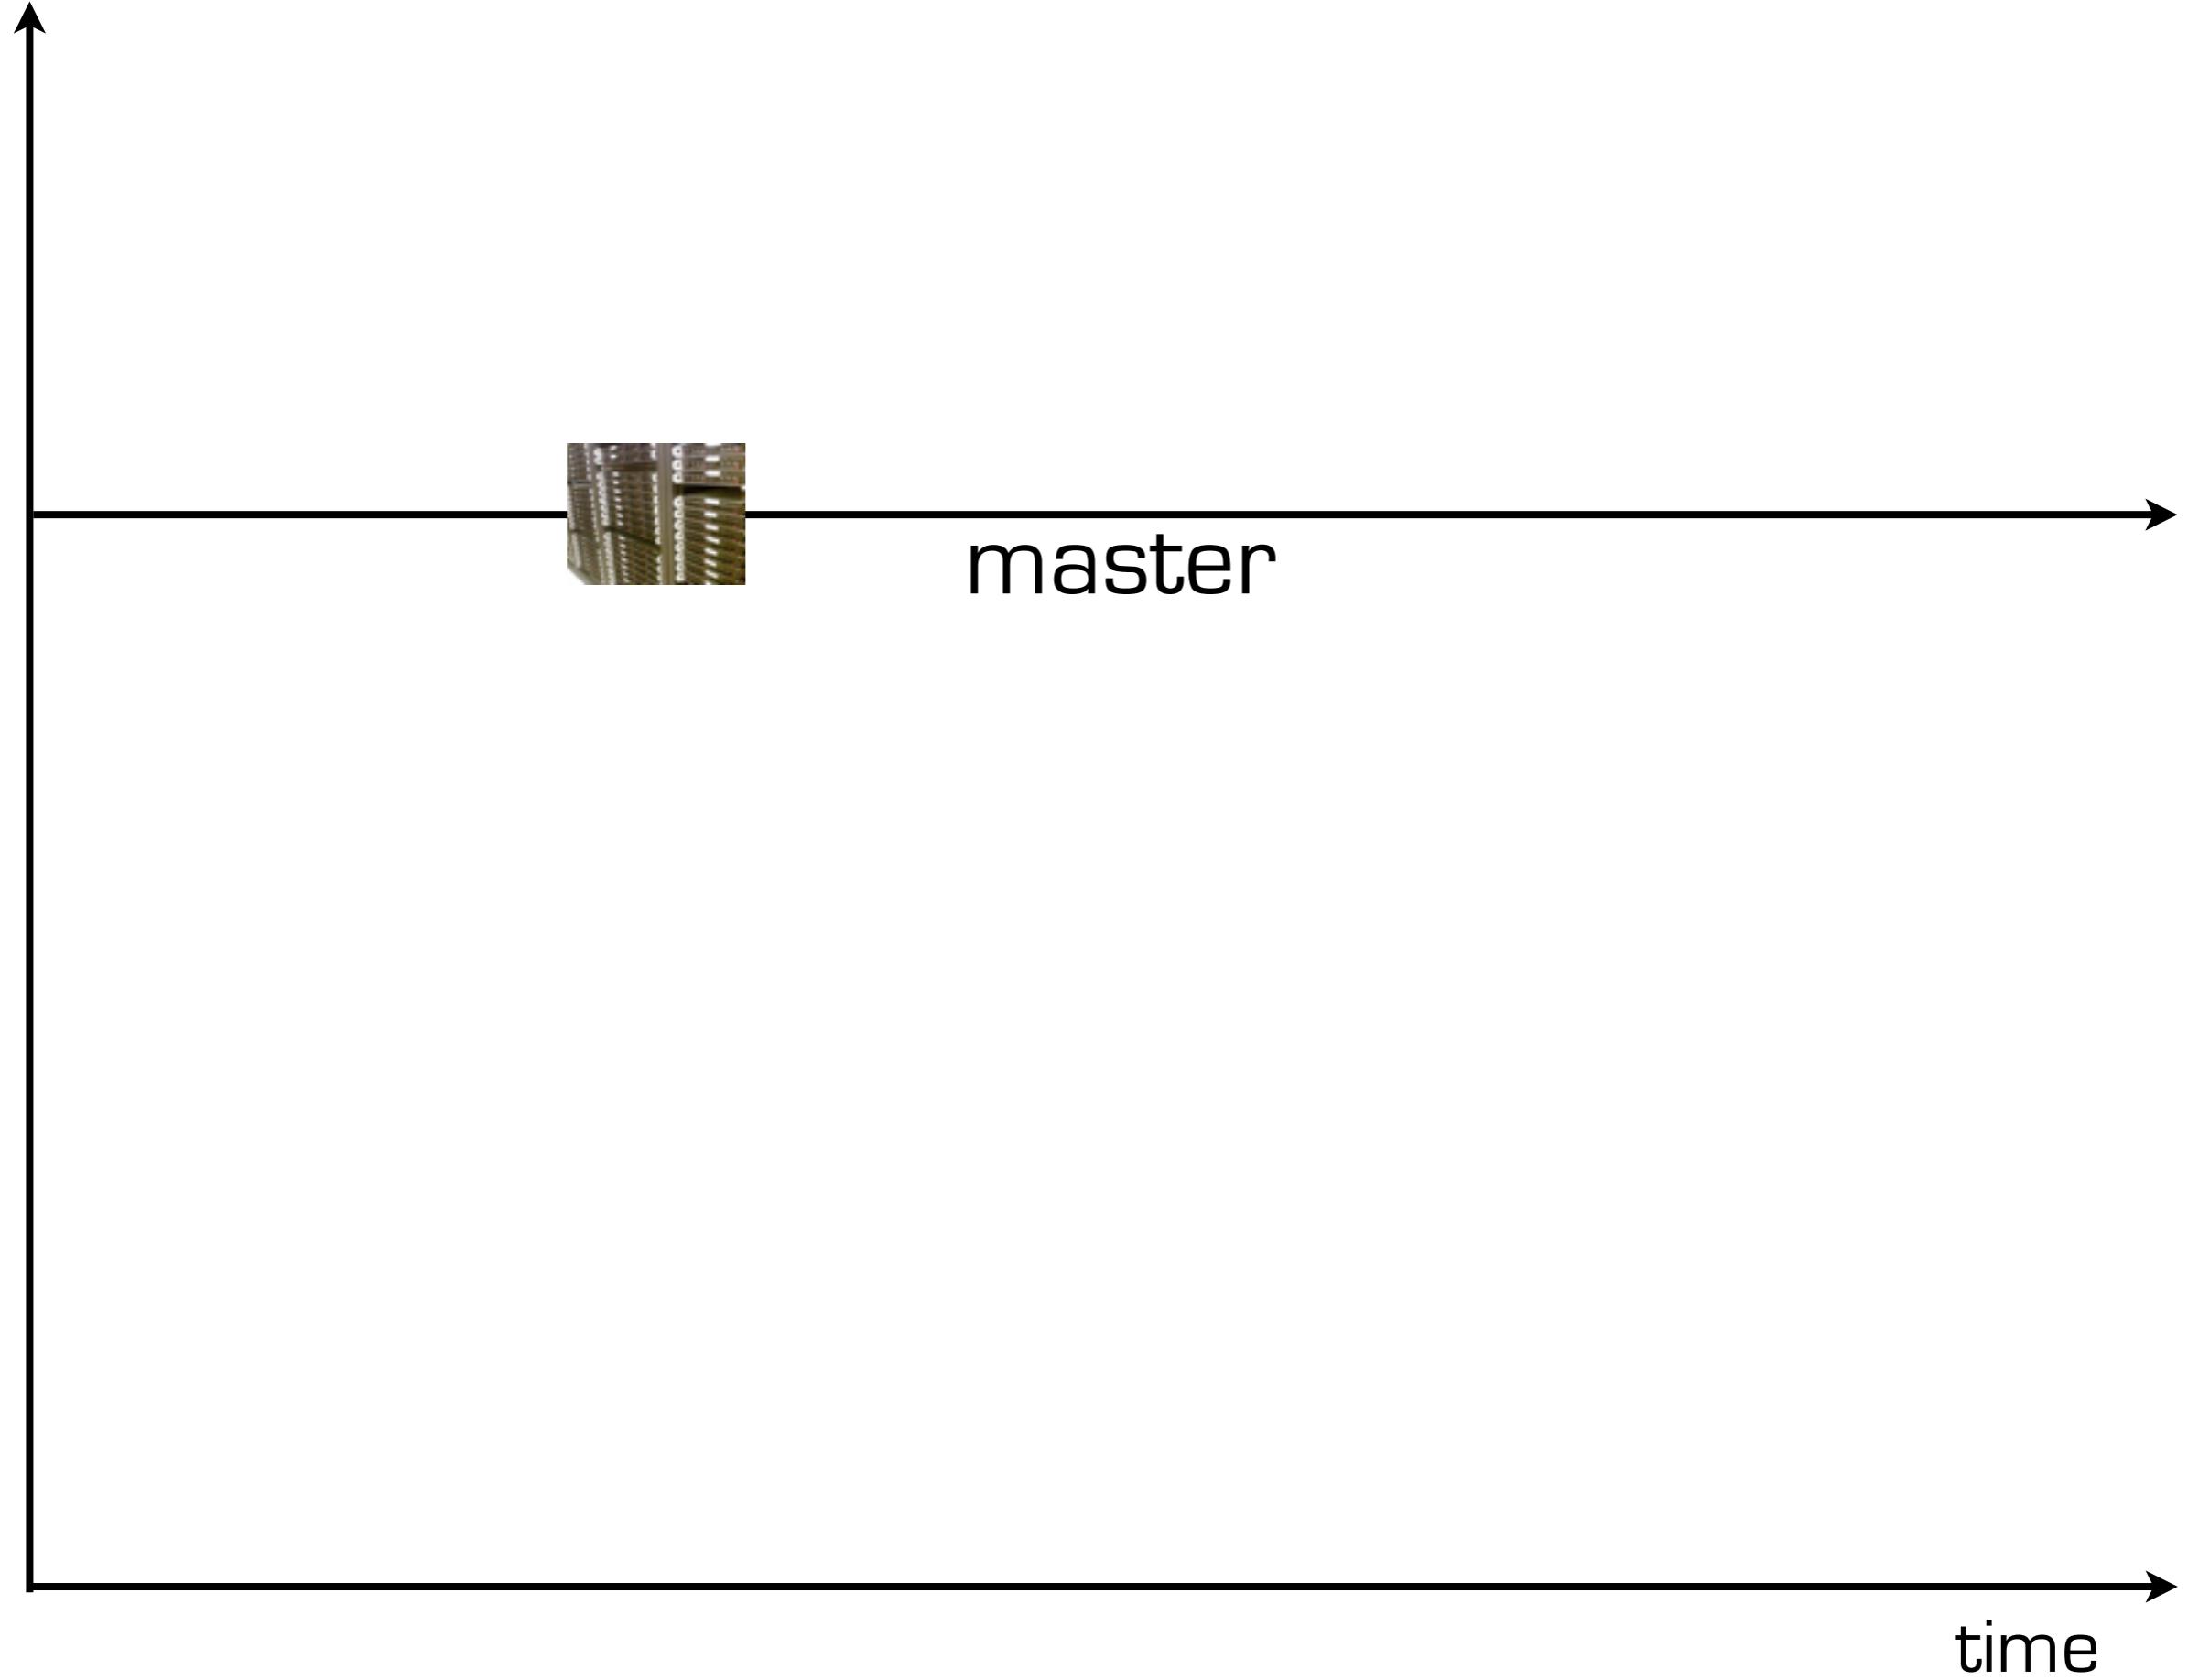


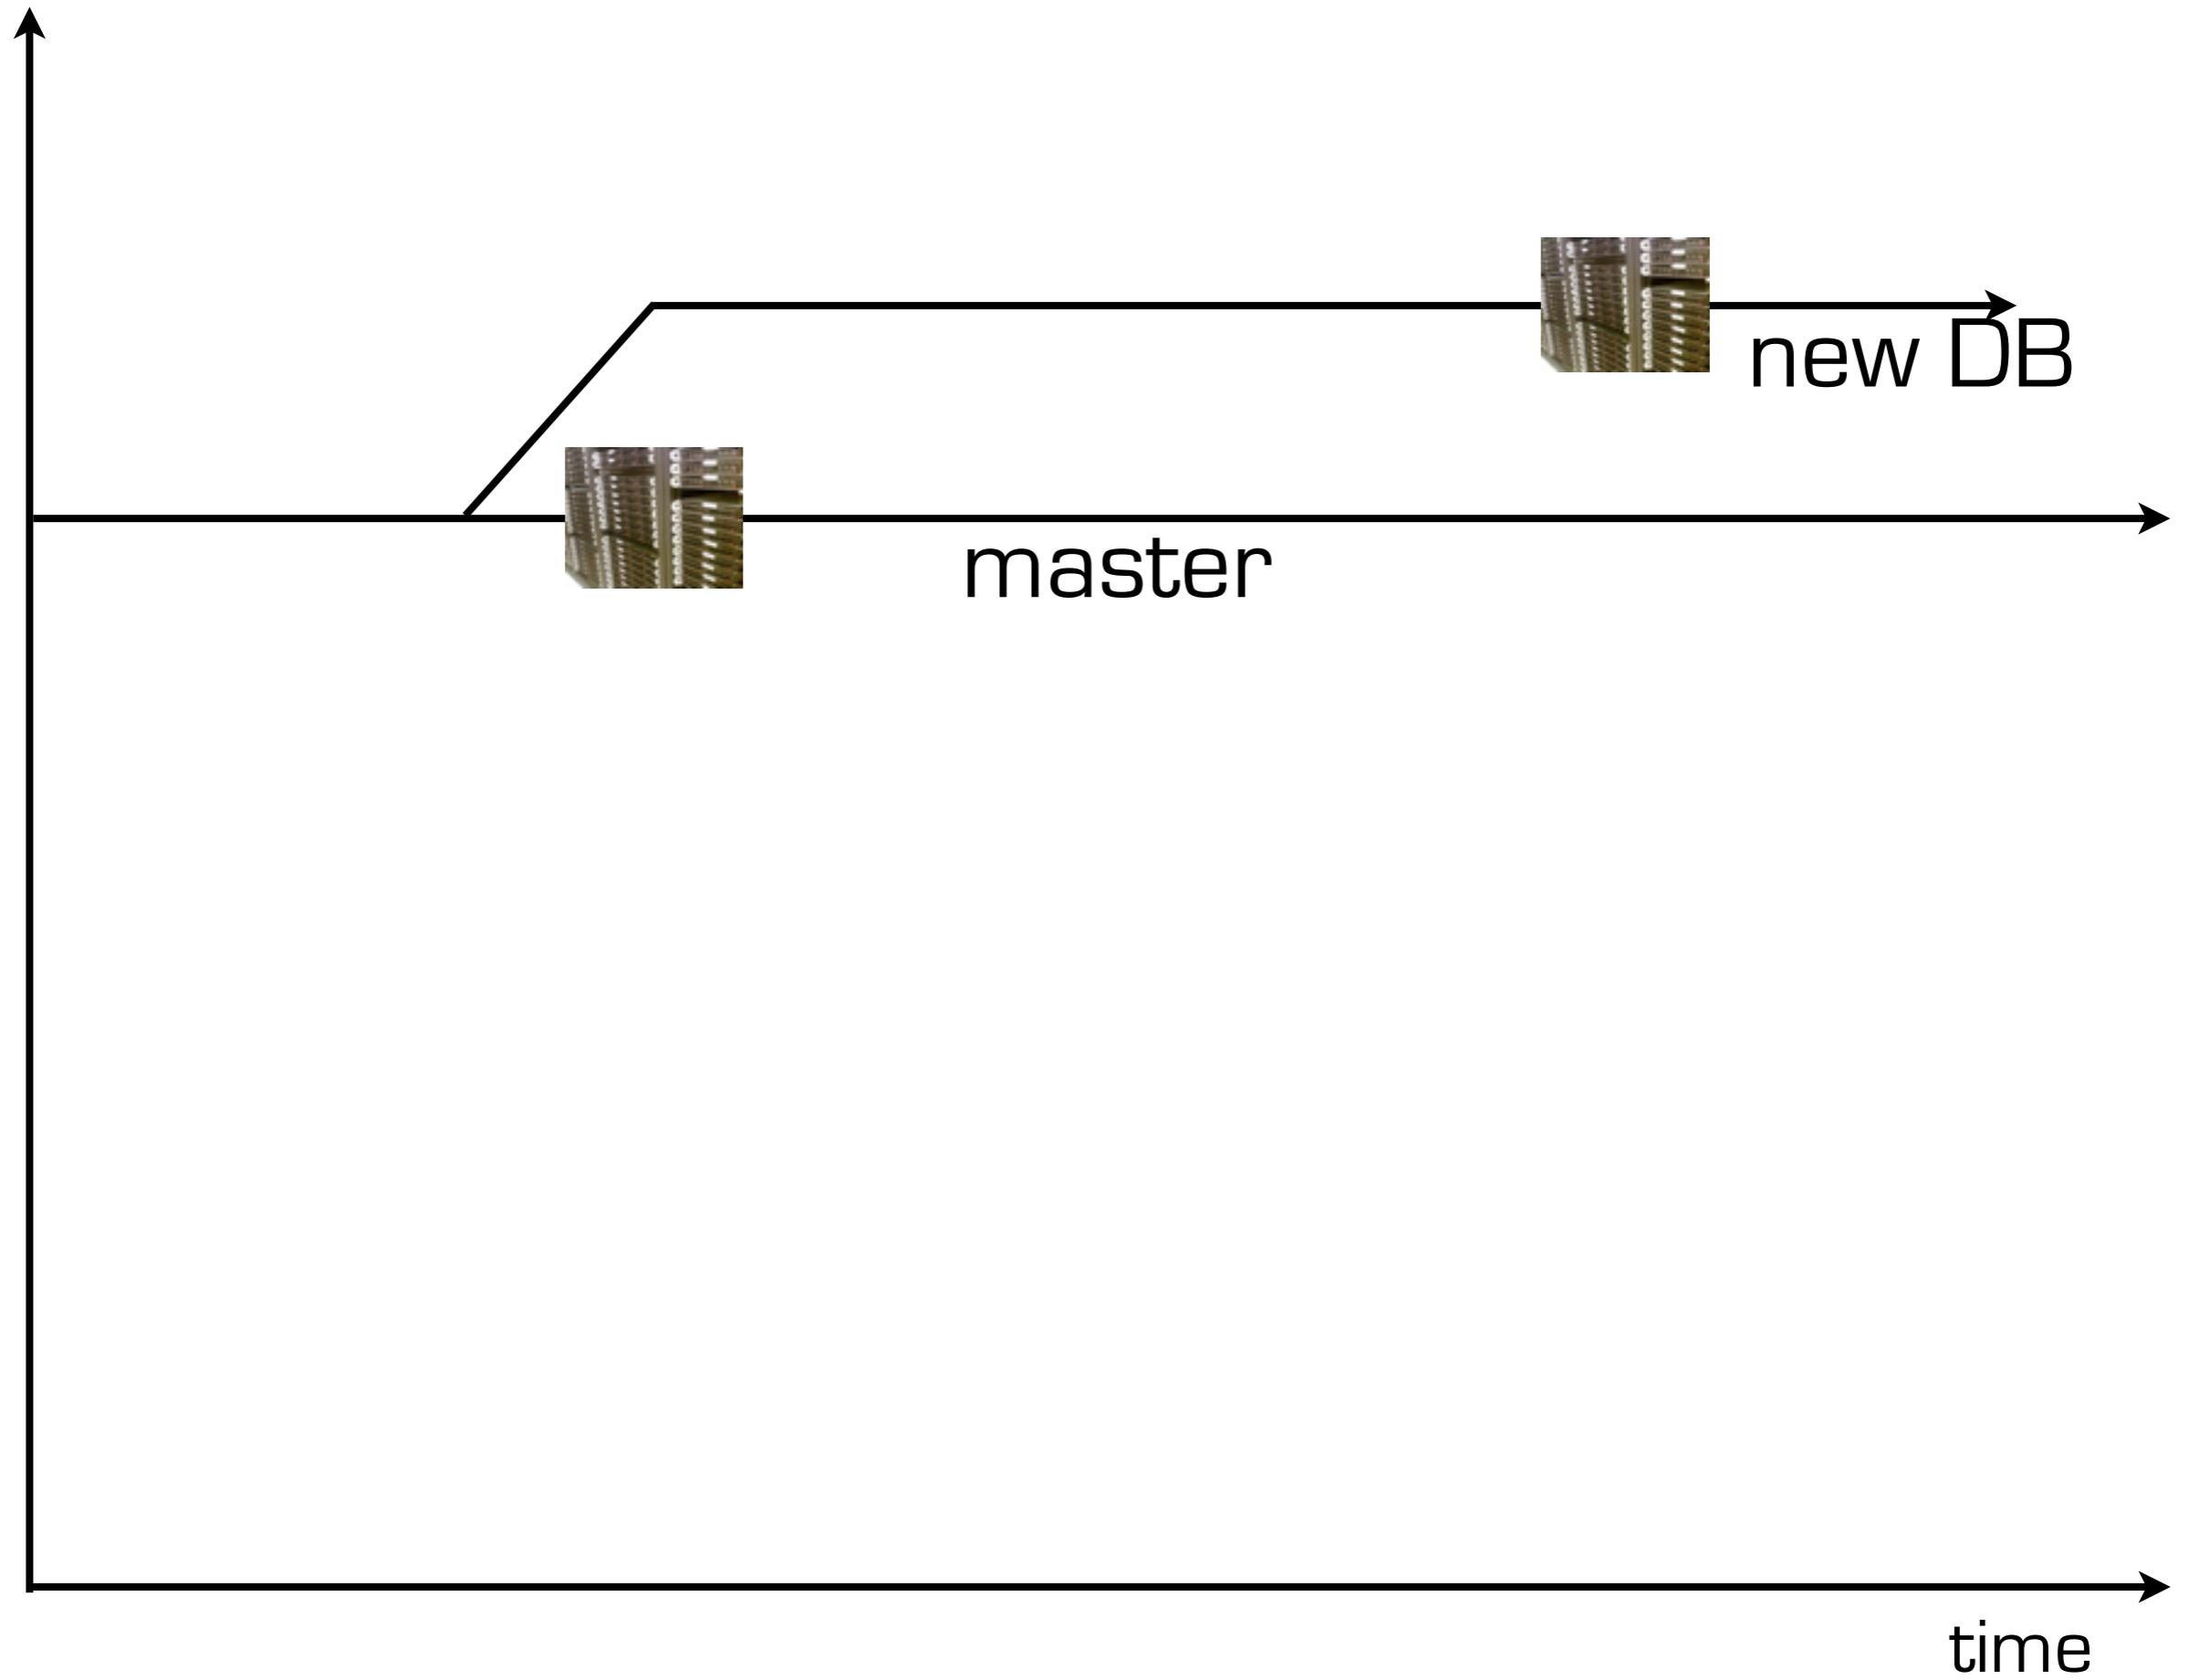
A diagram of a prism refracting light into a spectrum. A white line enters a triangular prism from the left, refracting into a spectrum of colors (red, orange, yellow, green, blue, purple) that exit the prism to the right.

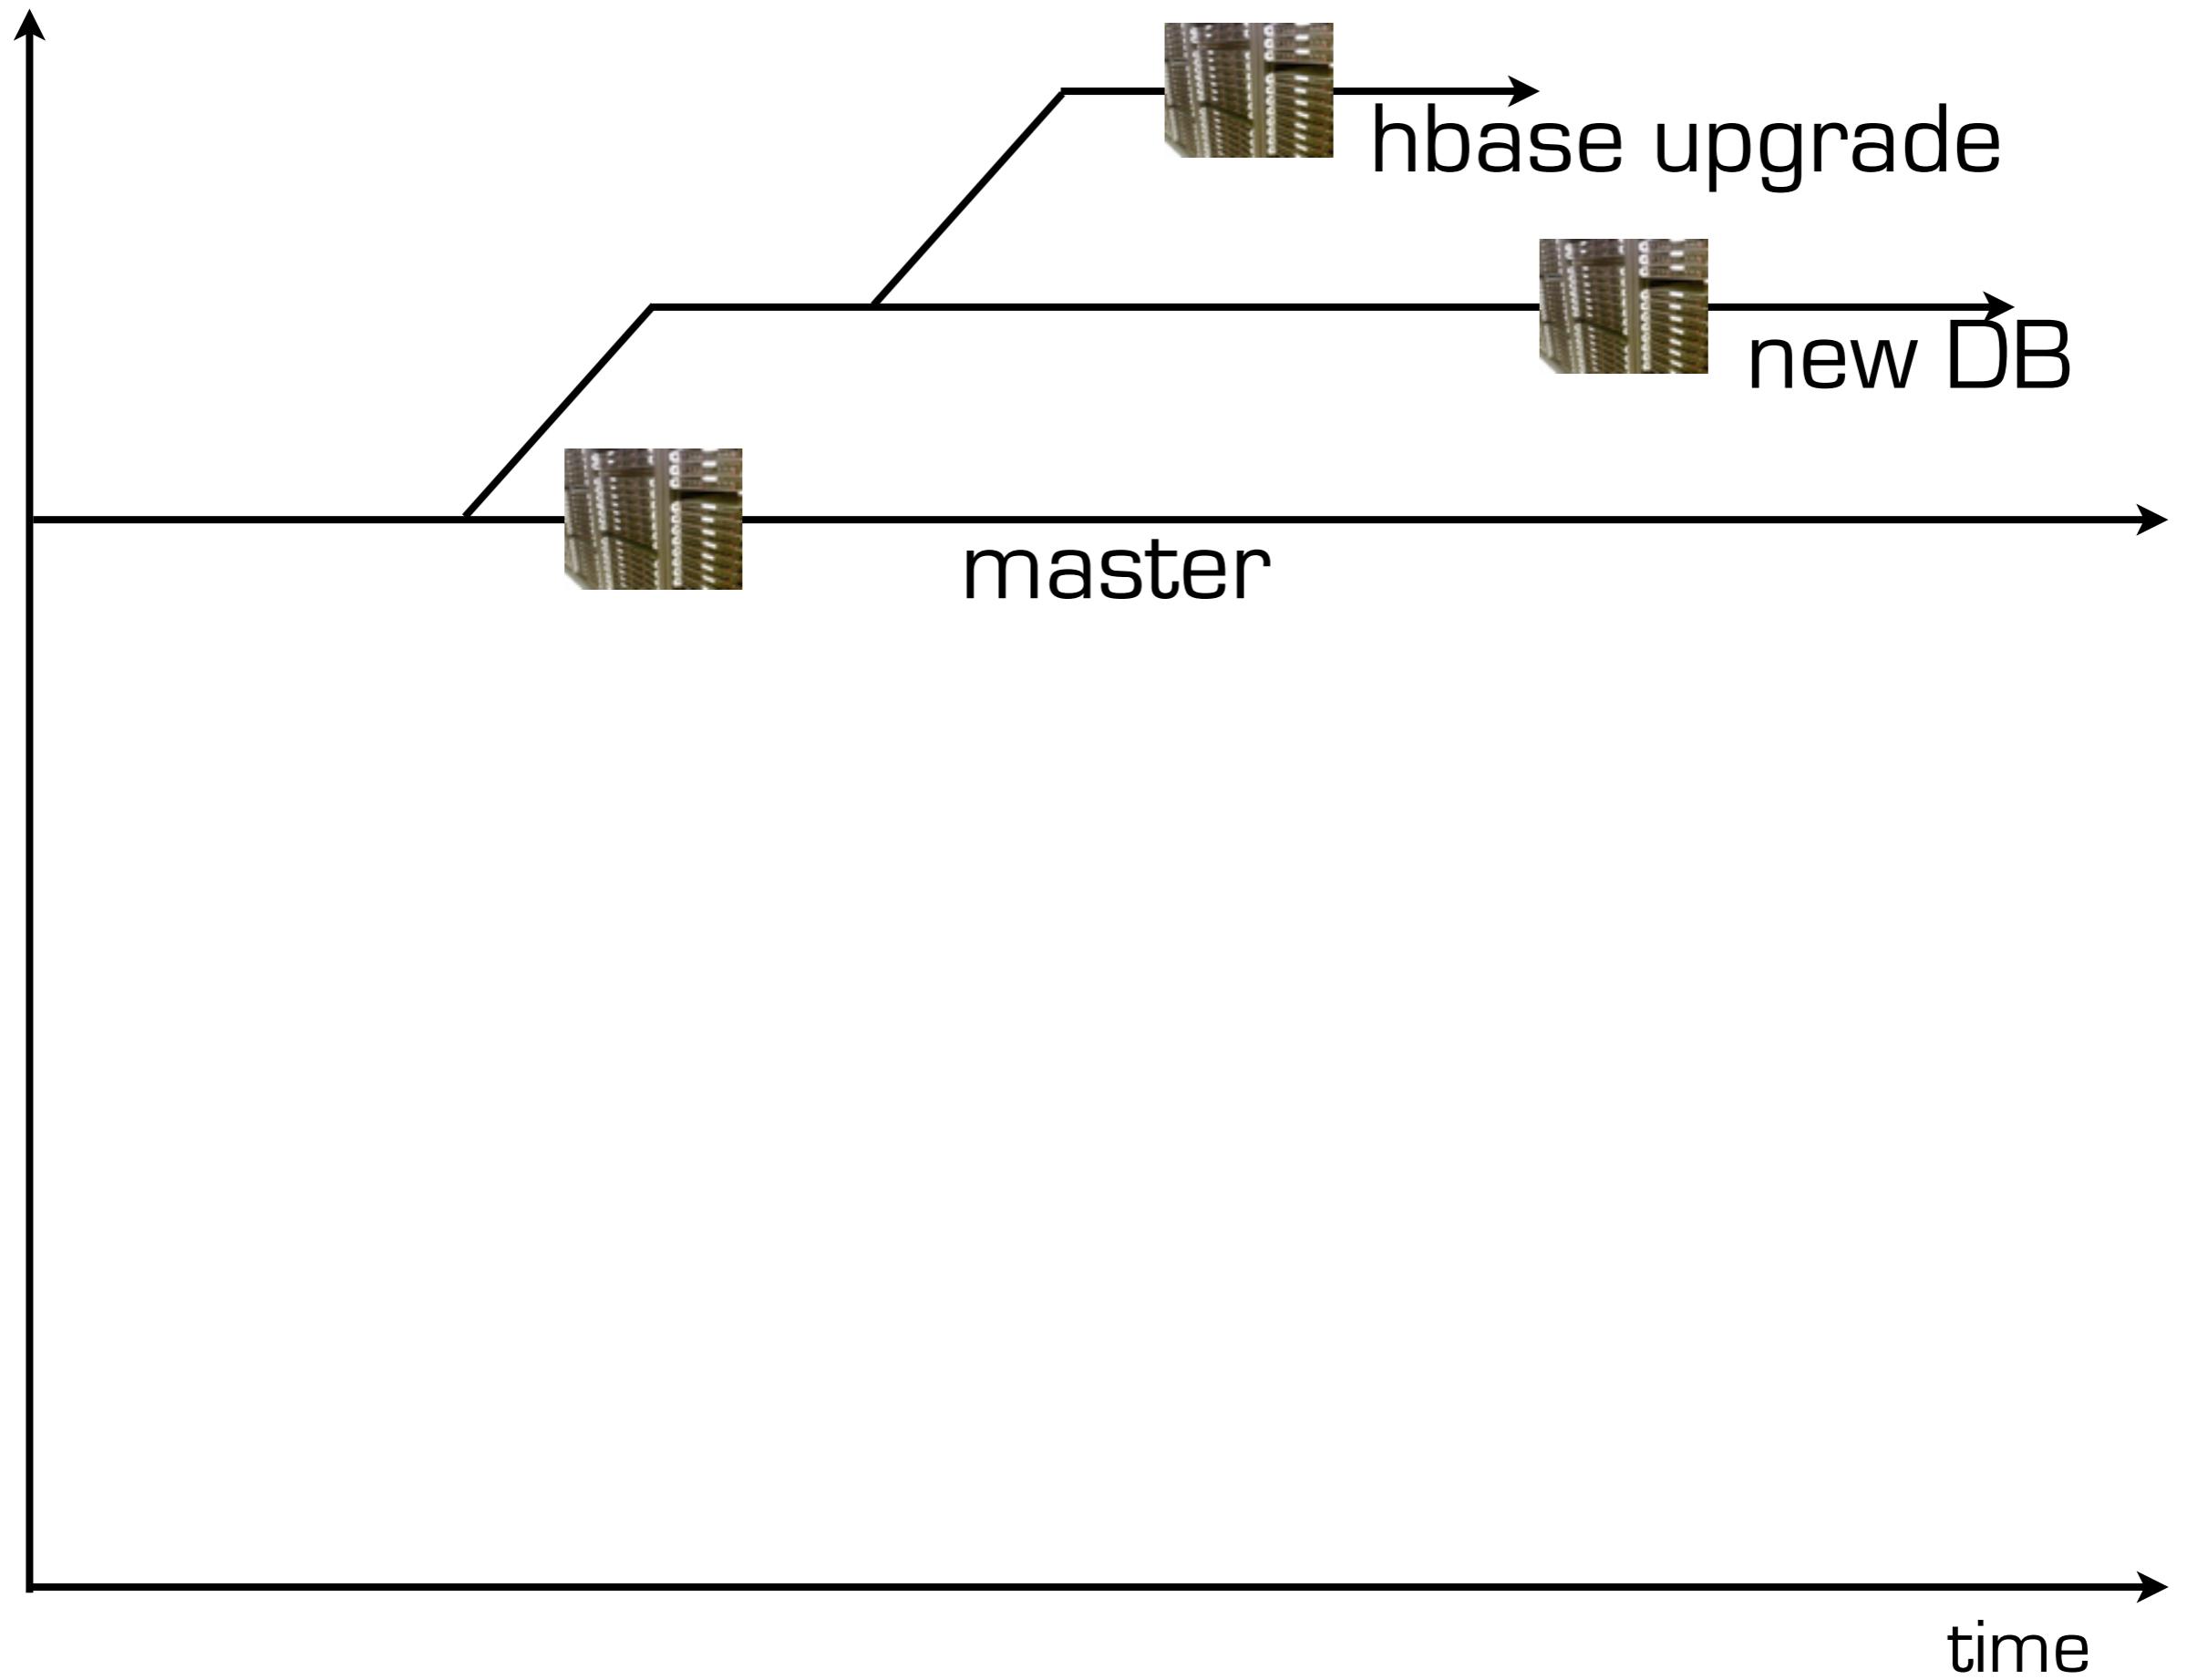
time

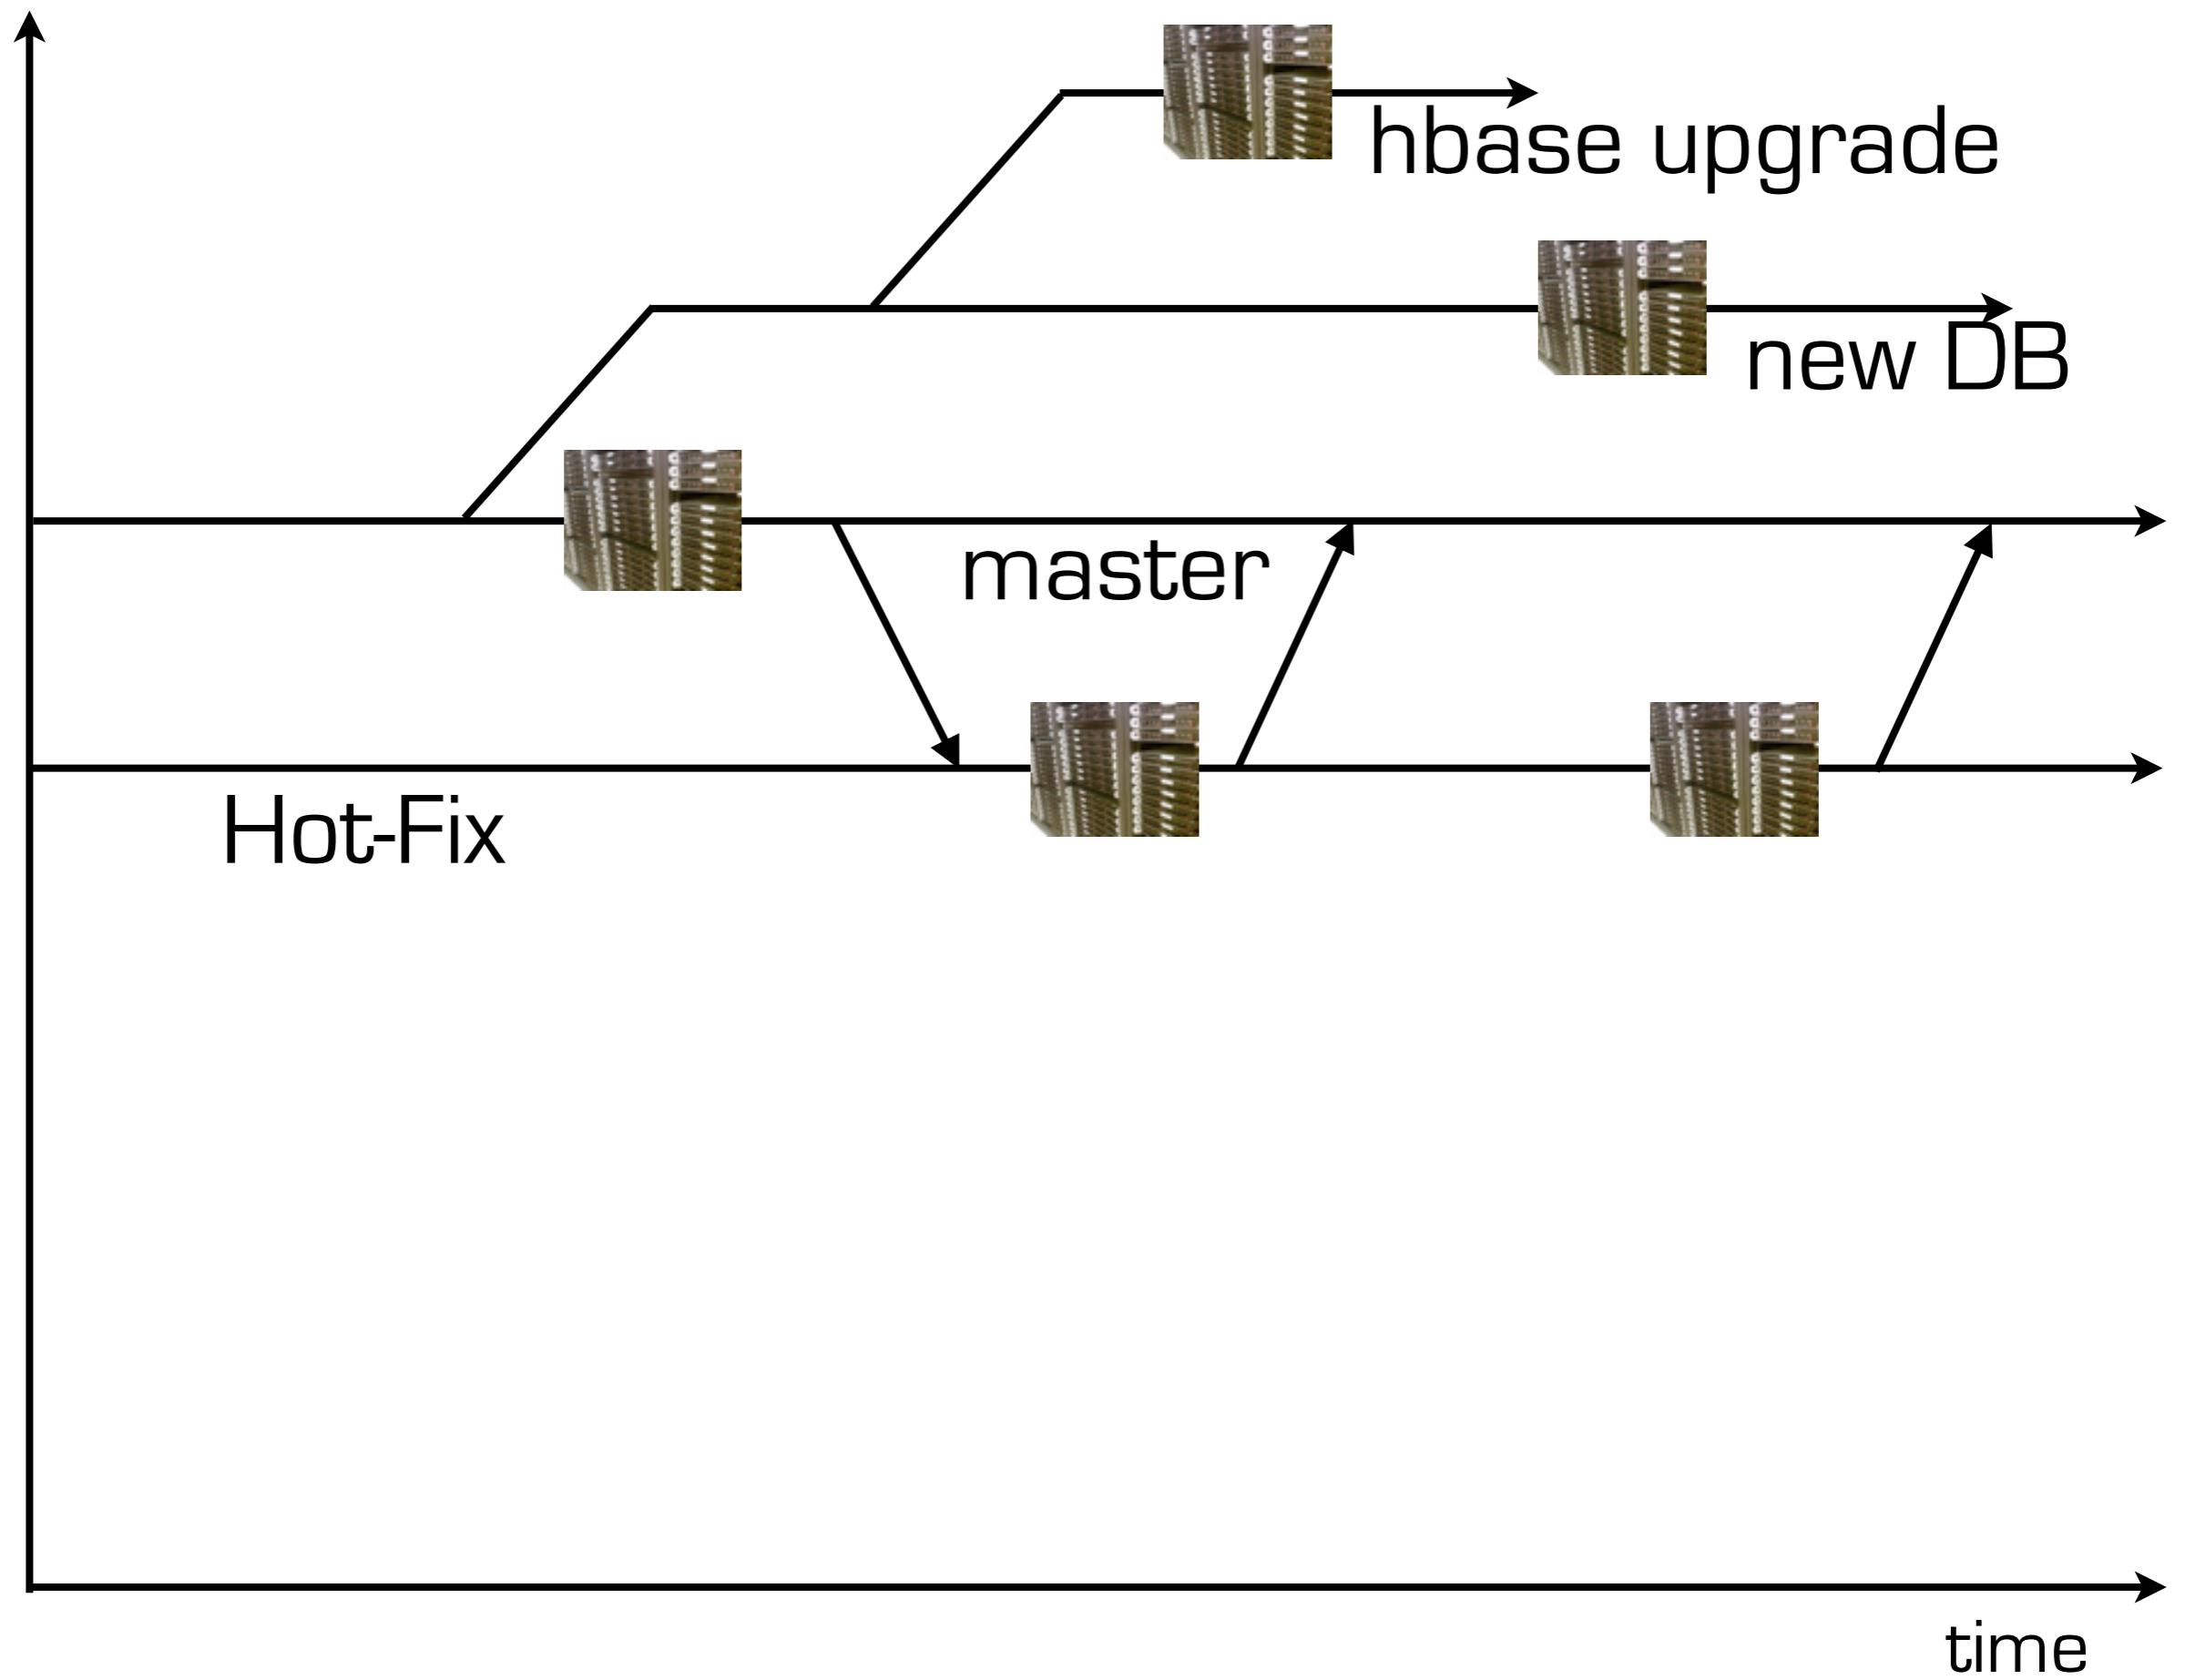


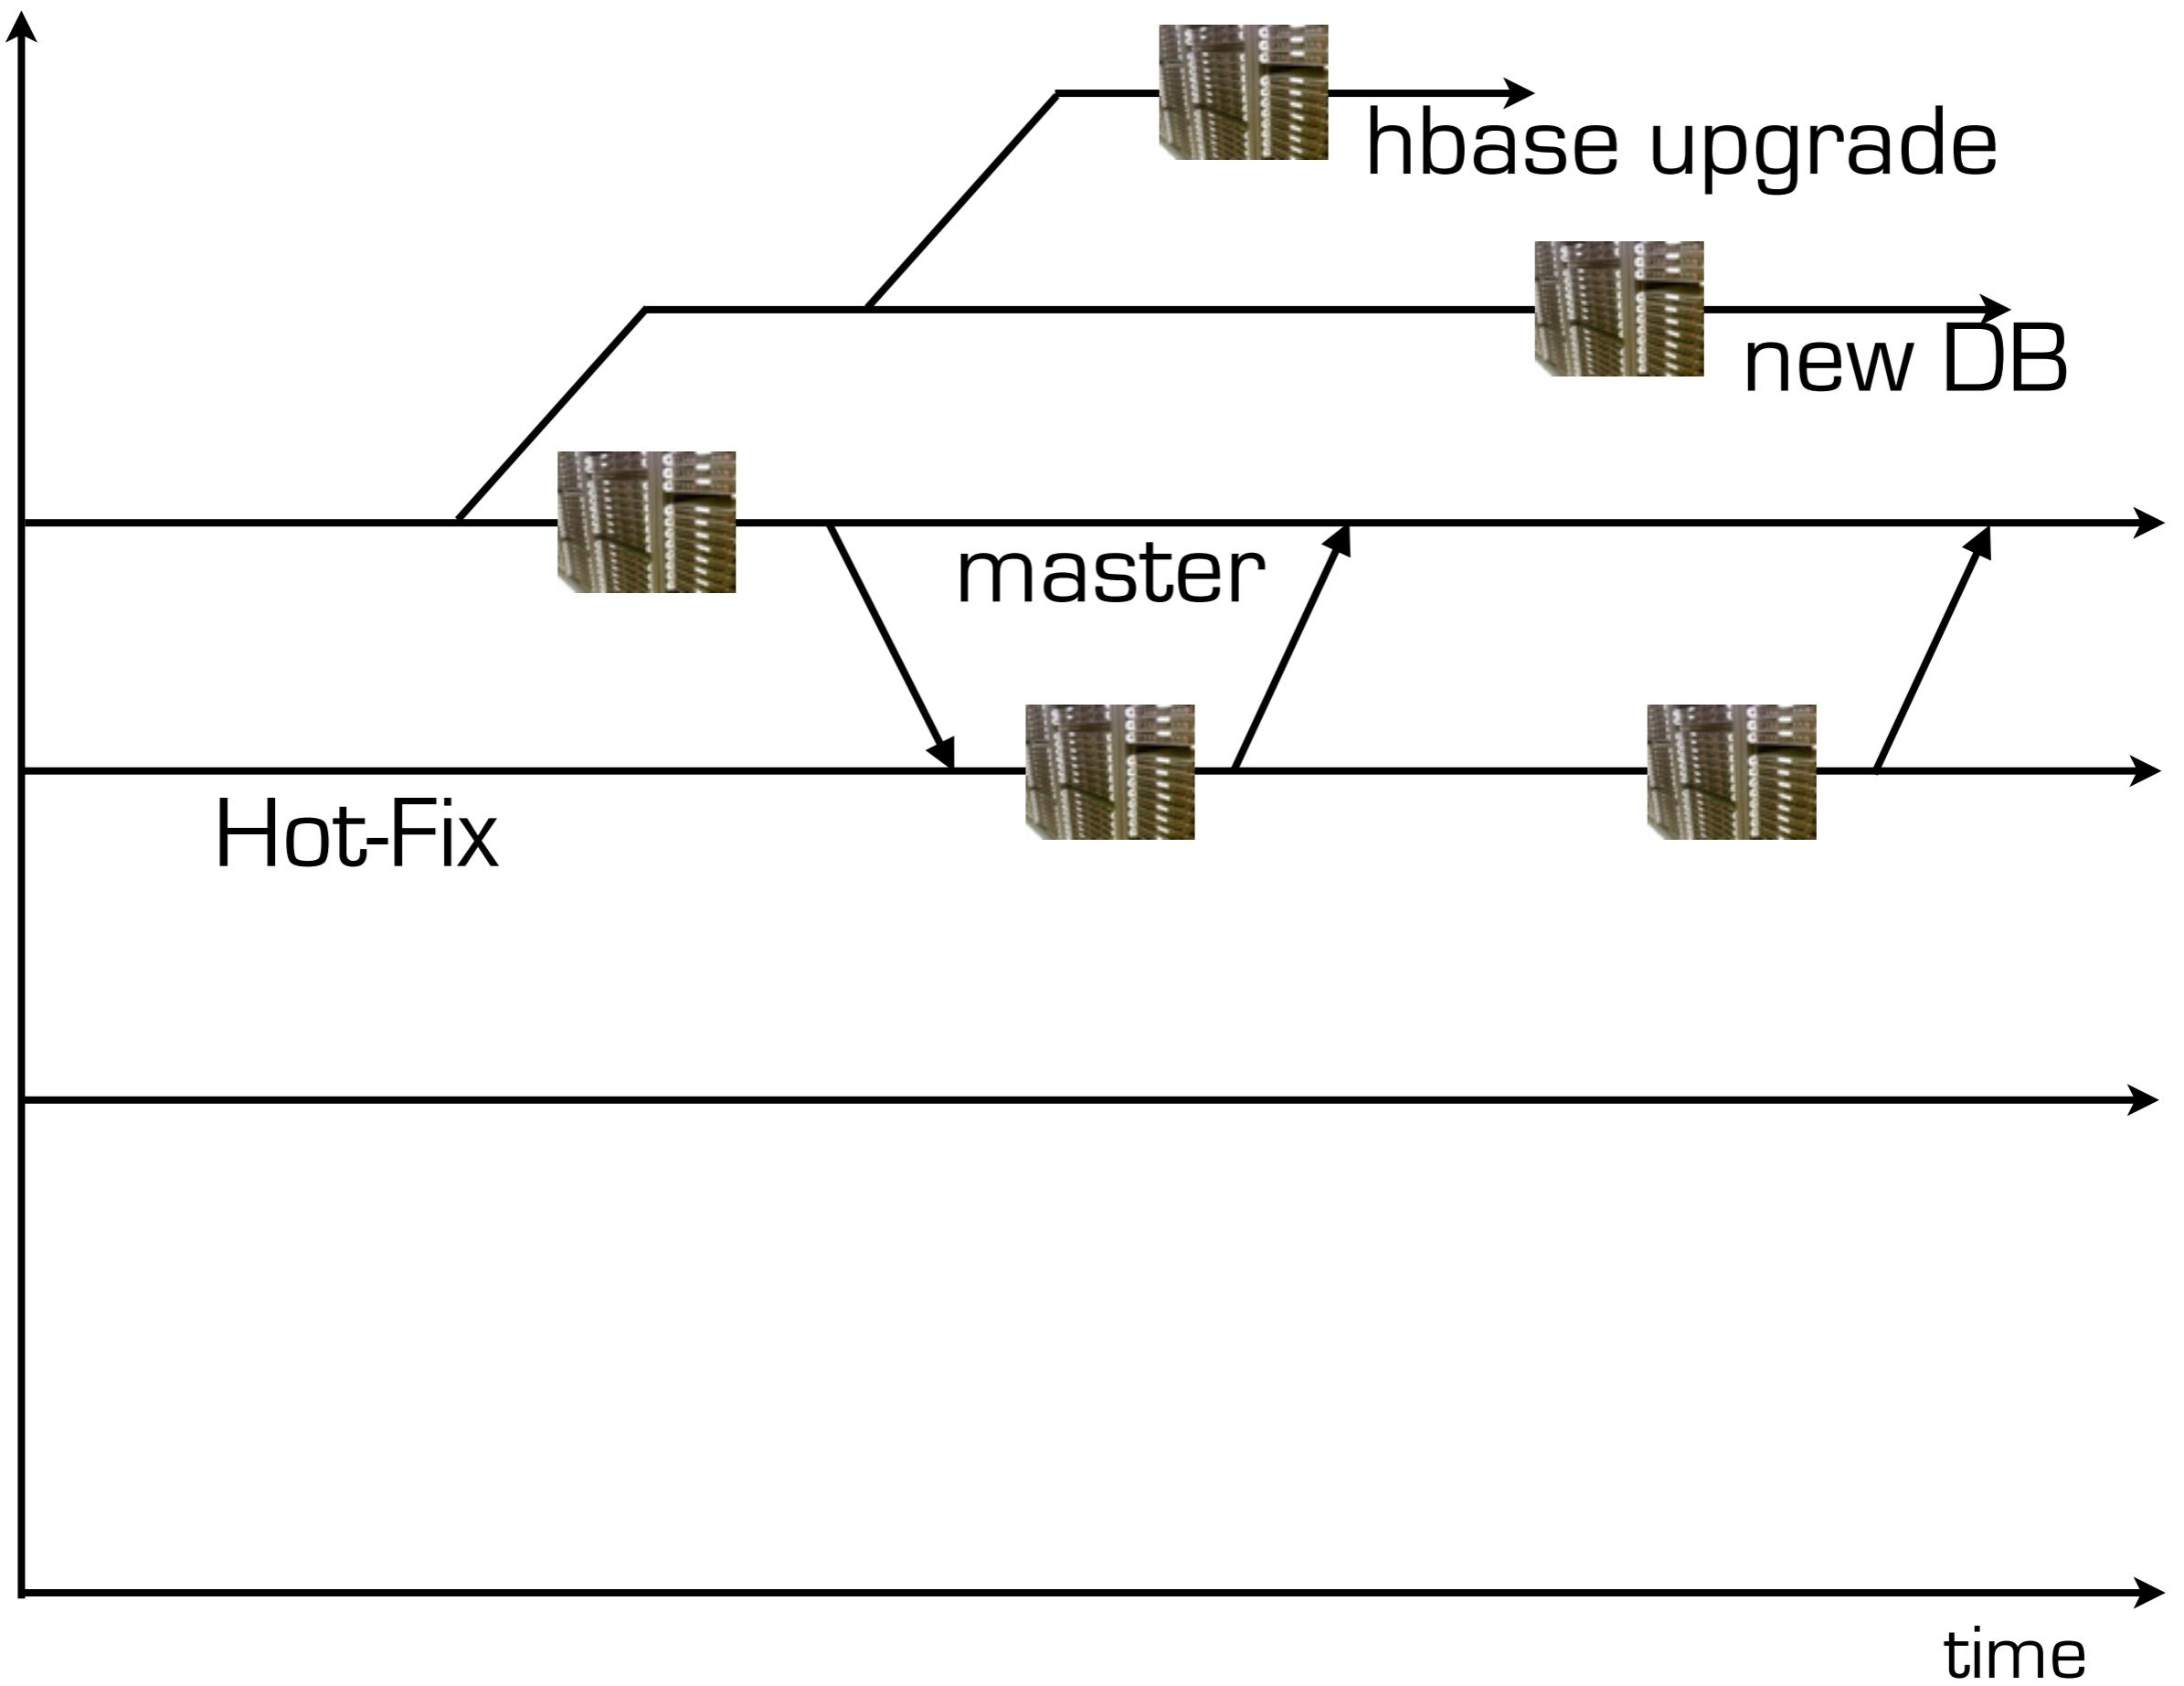
access access access access access

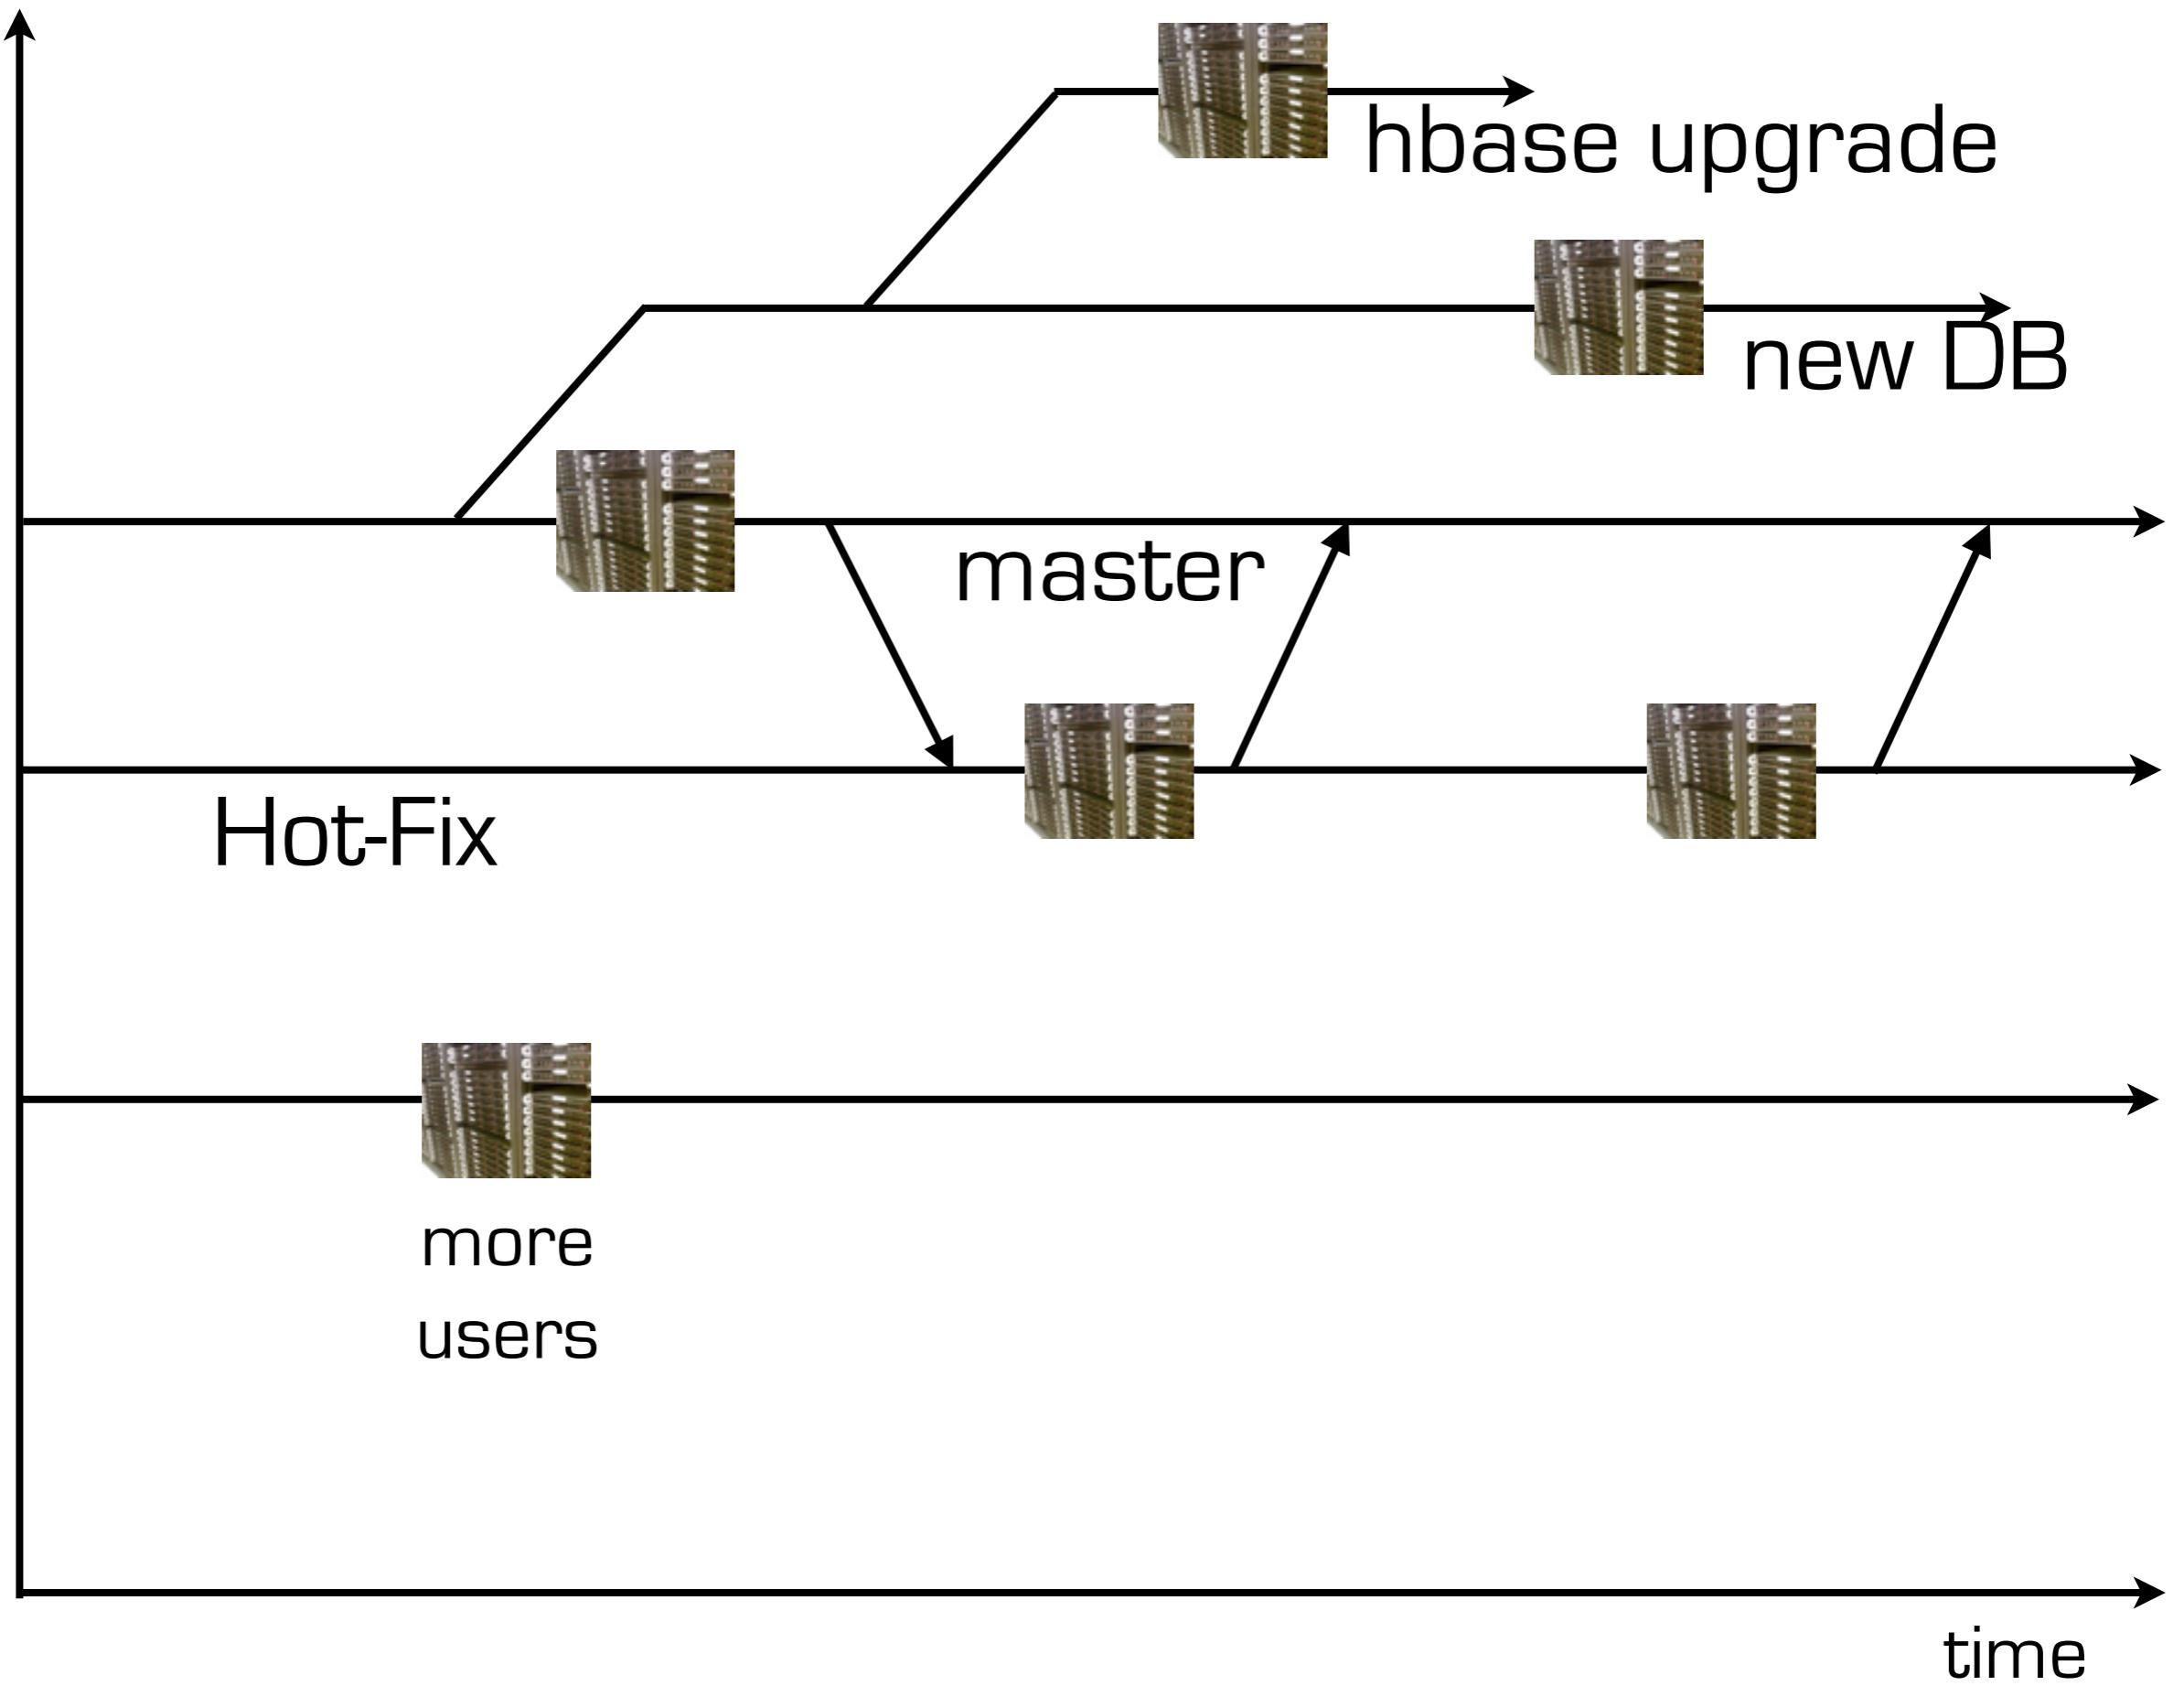


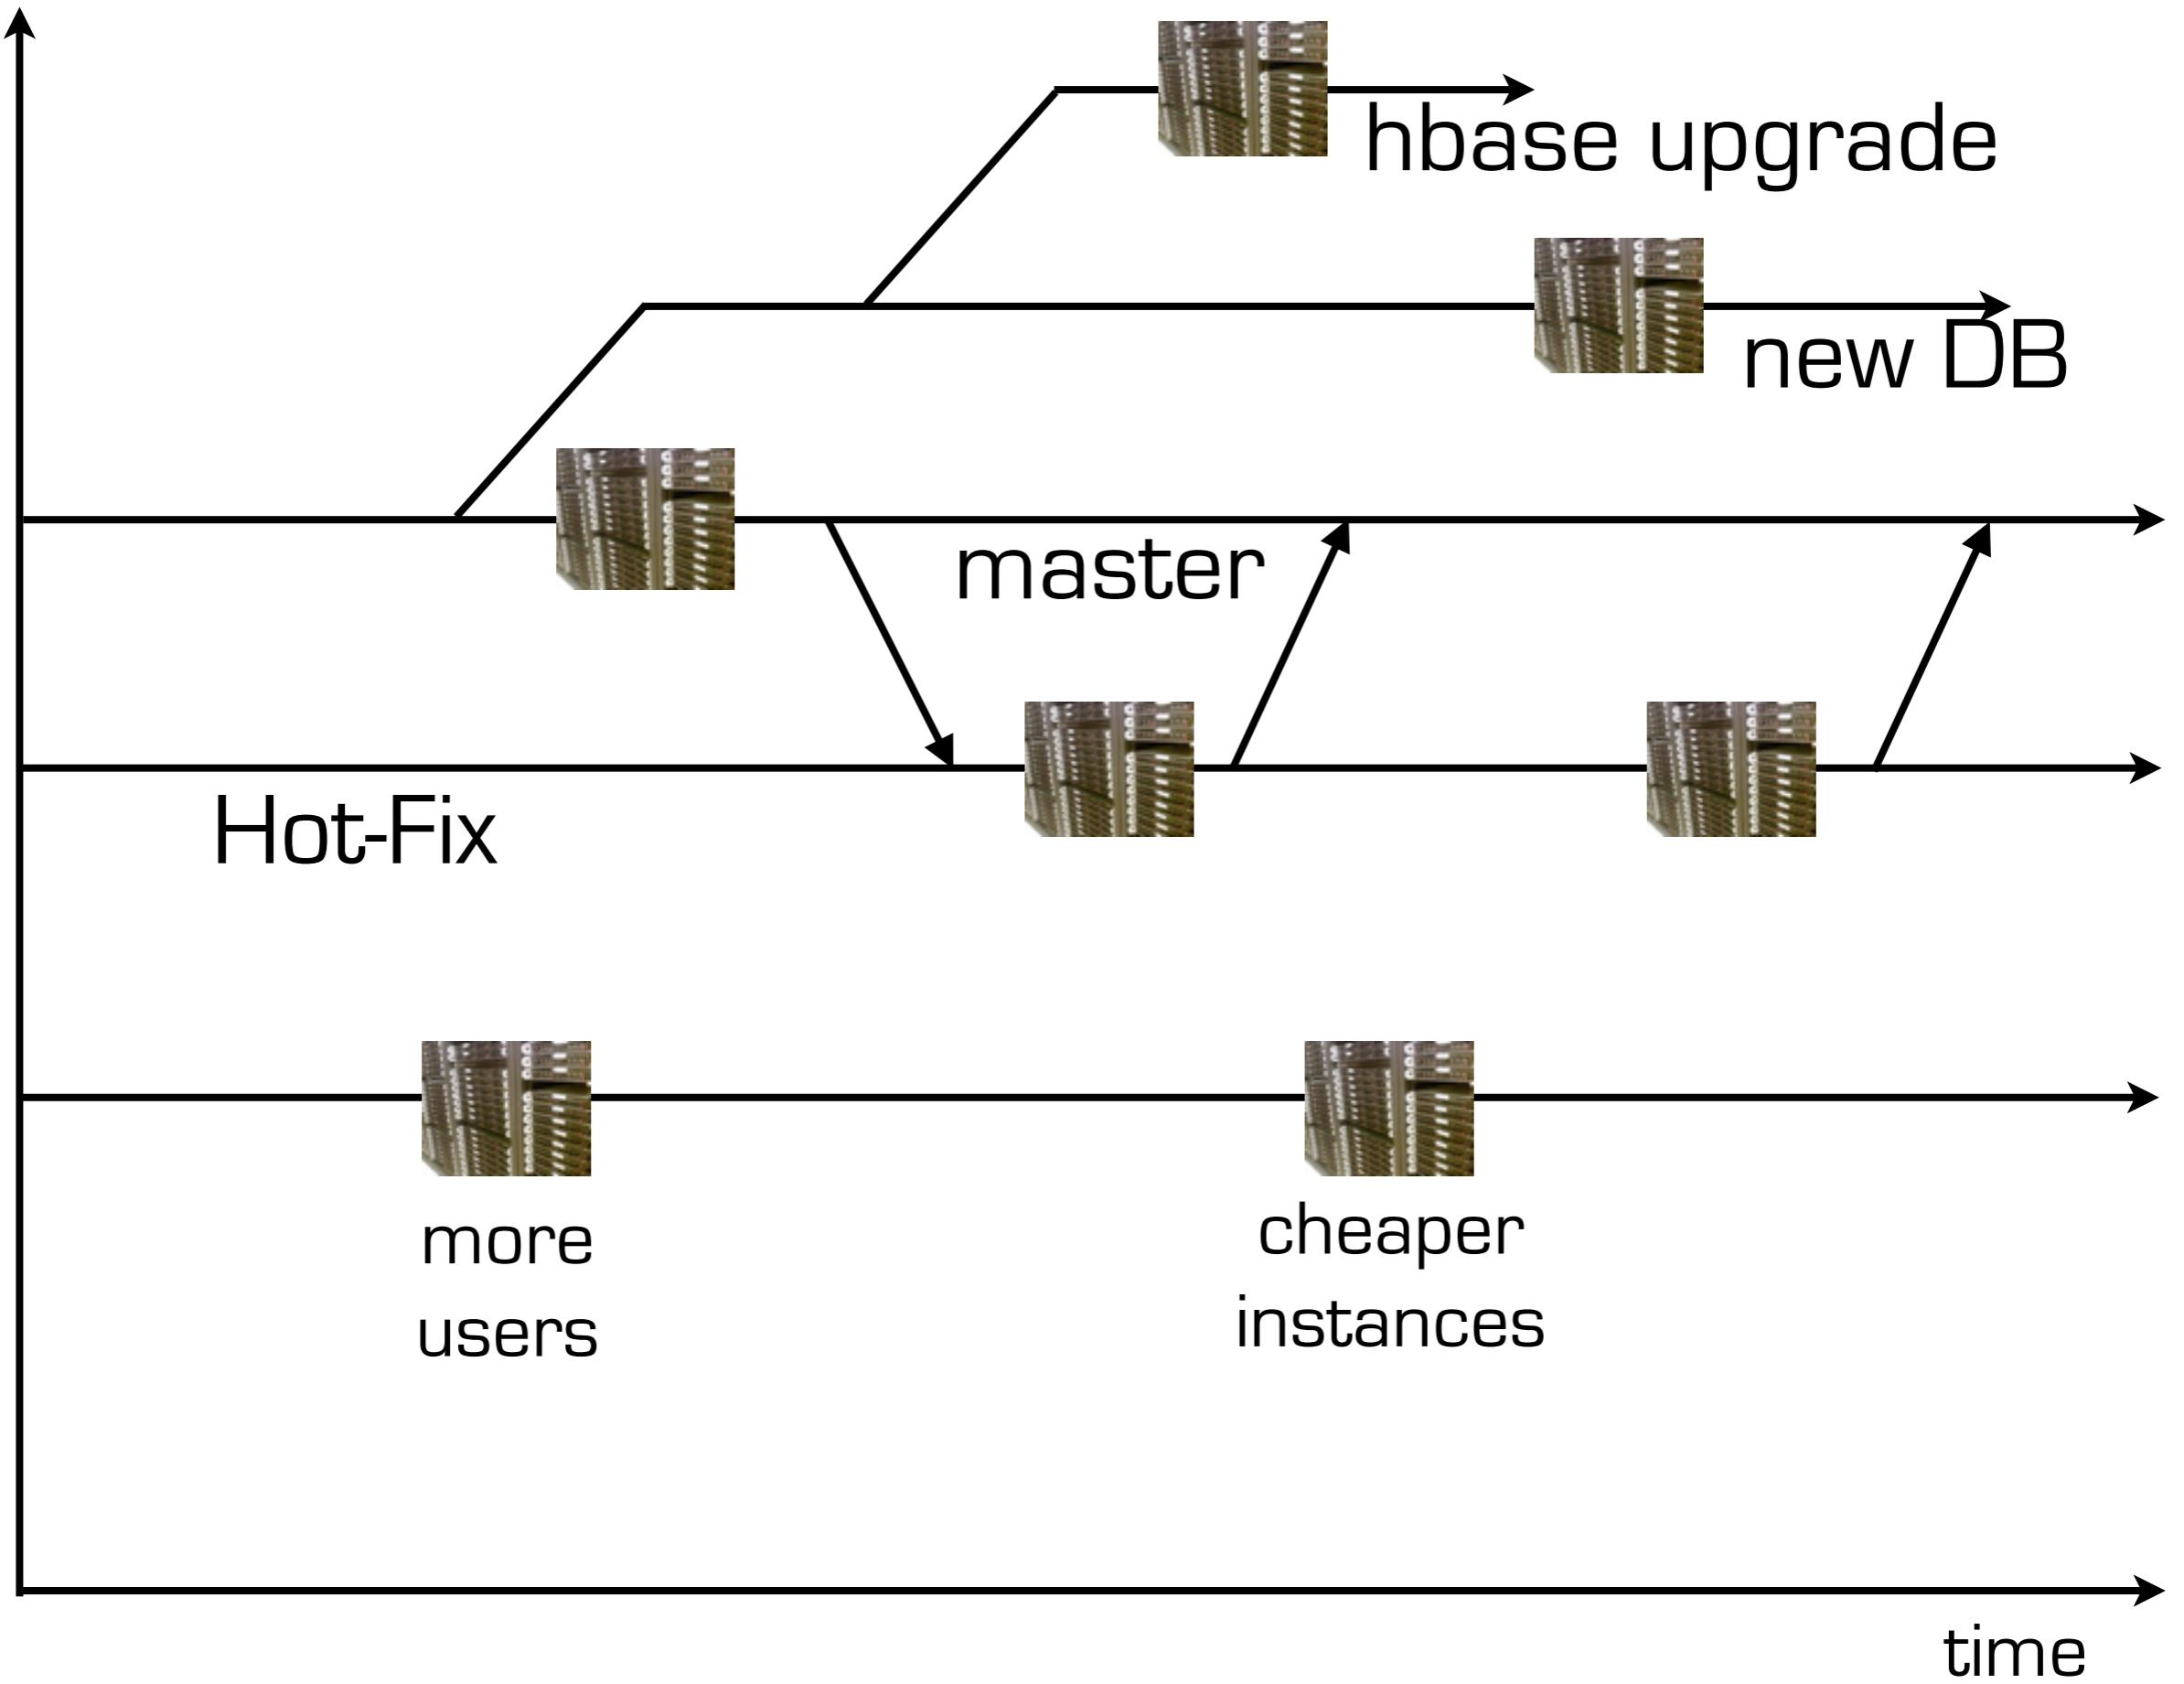


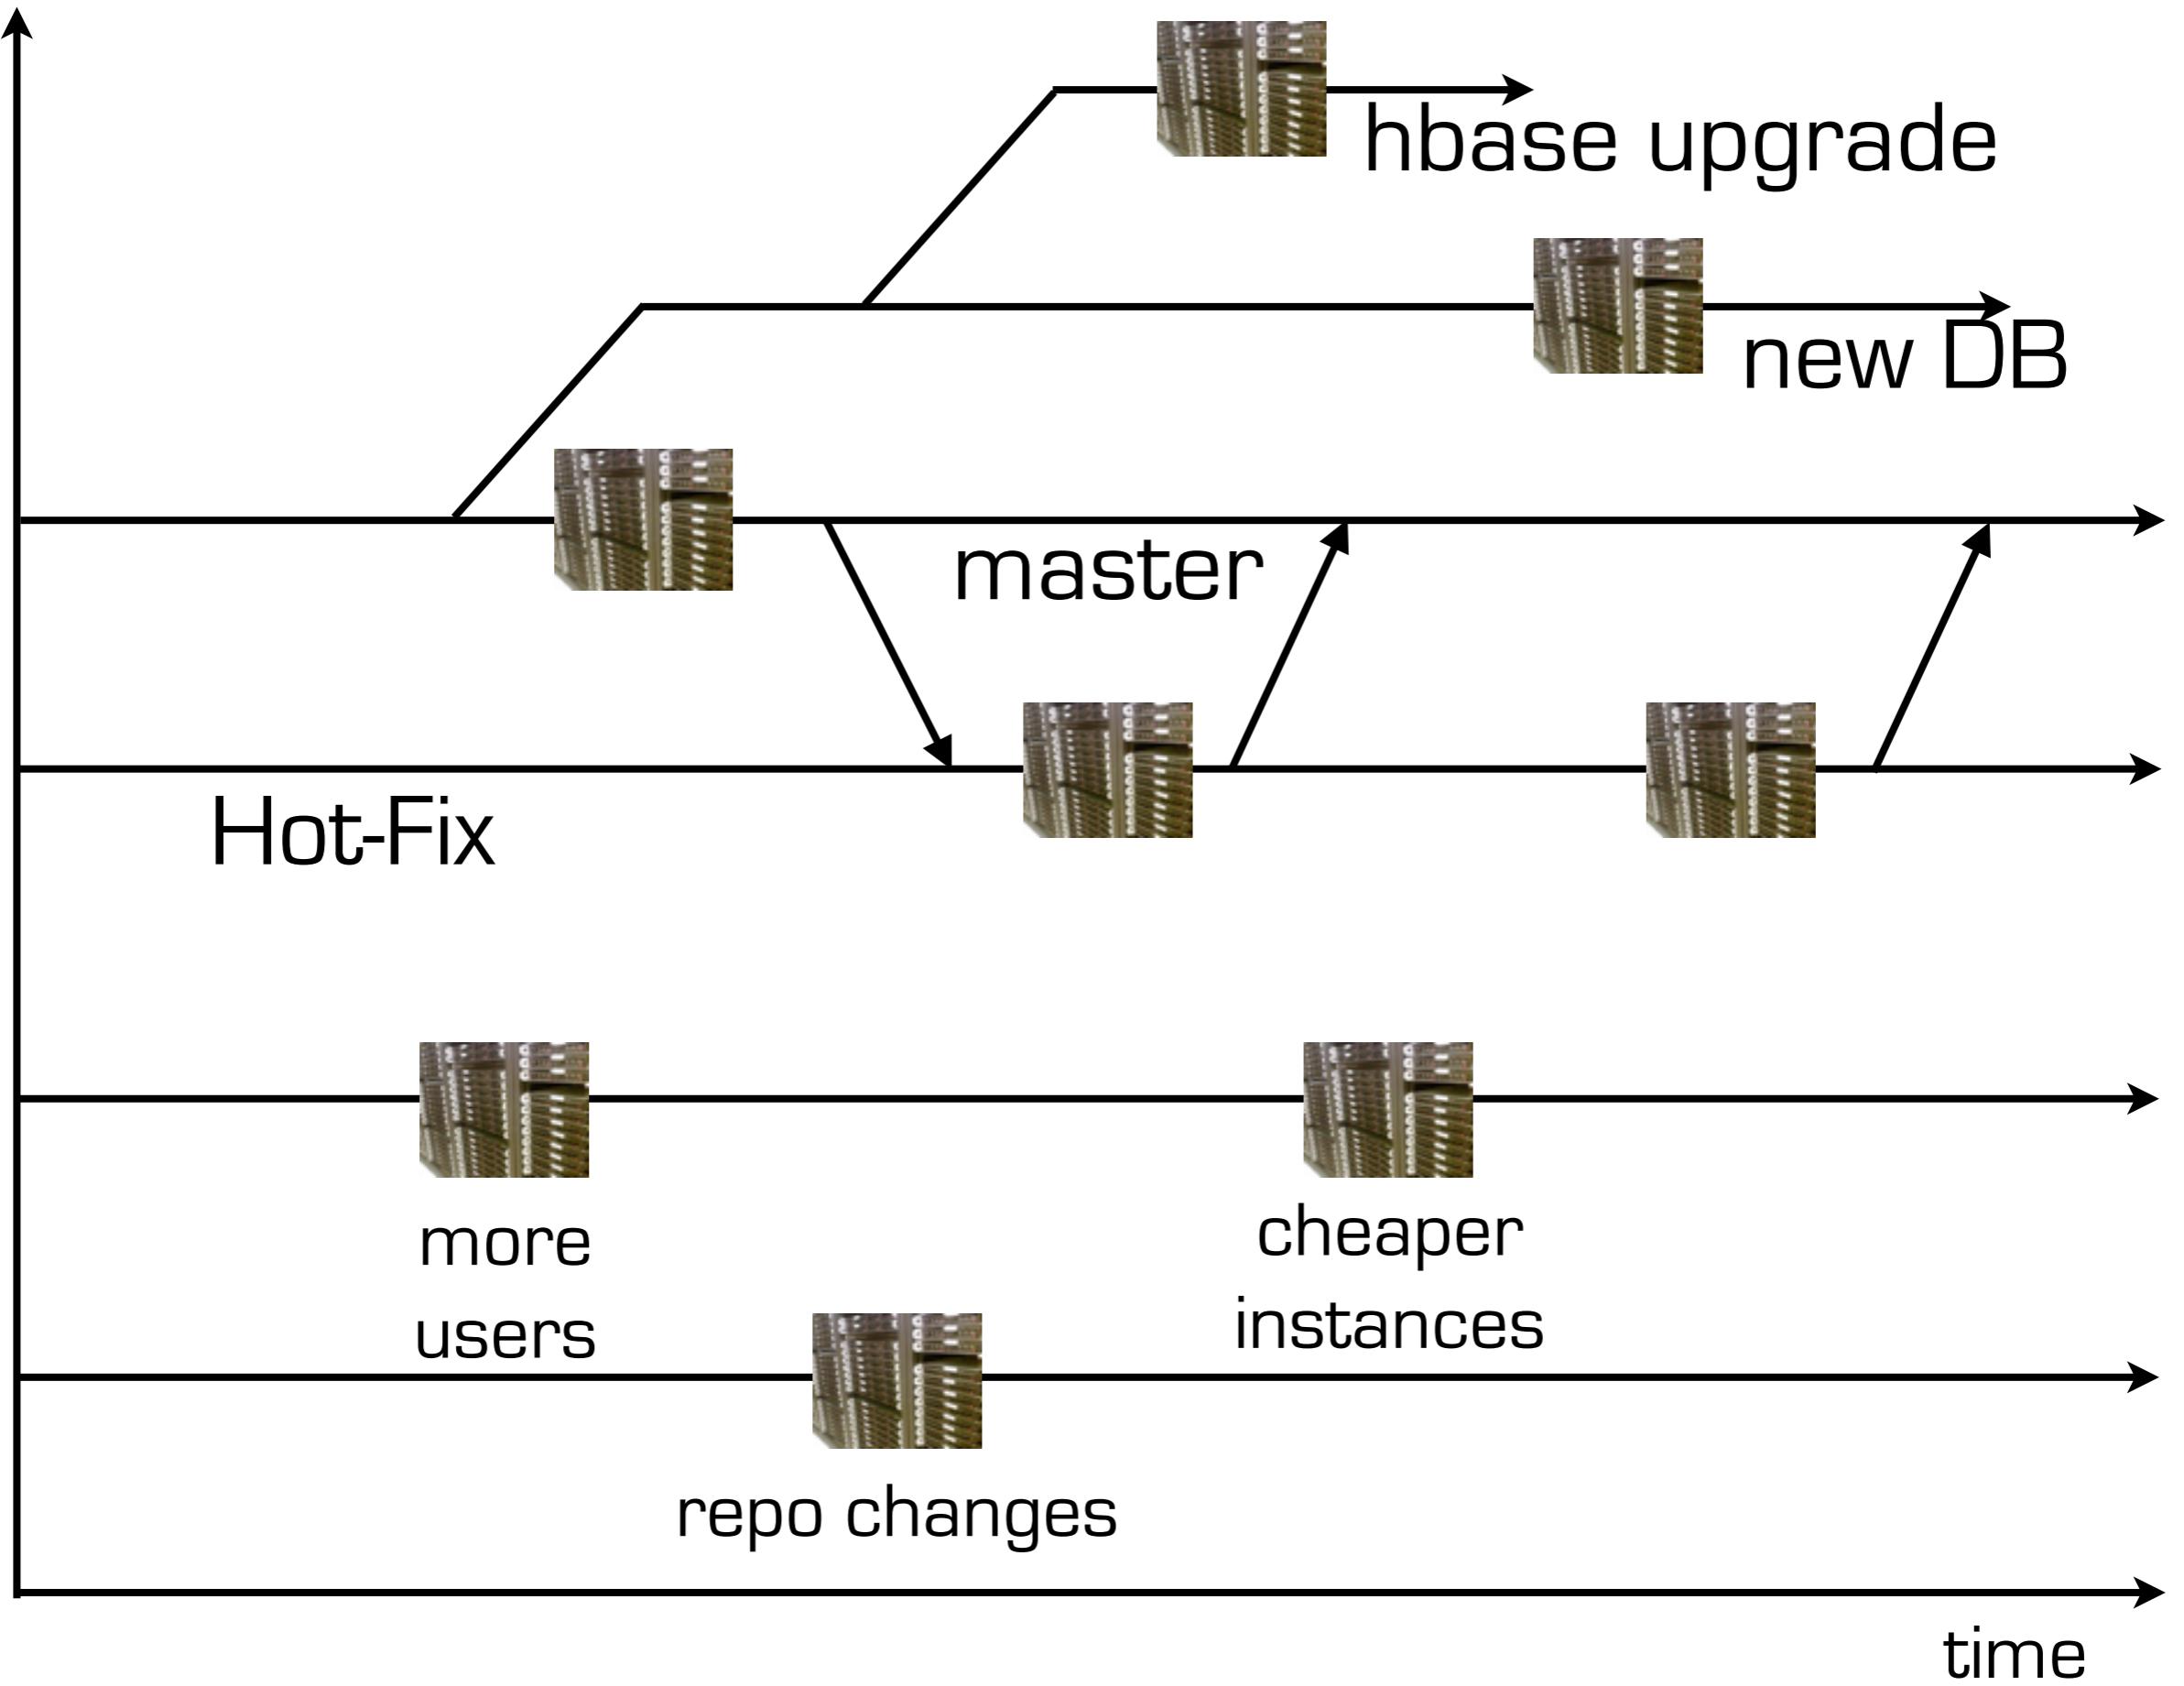








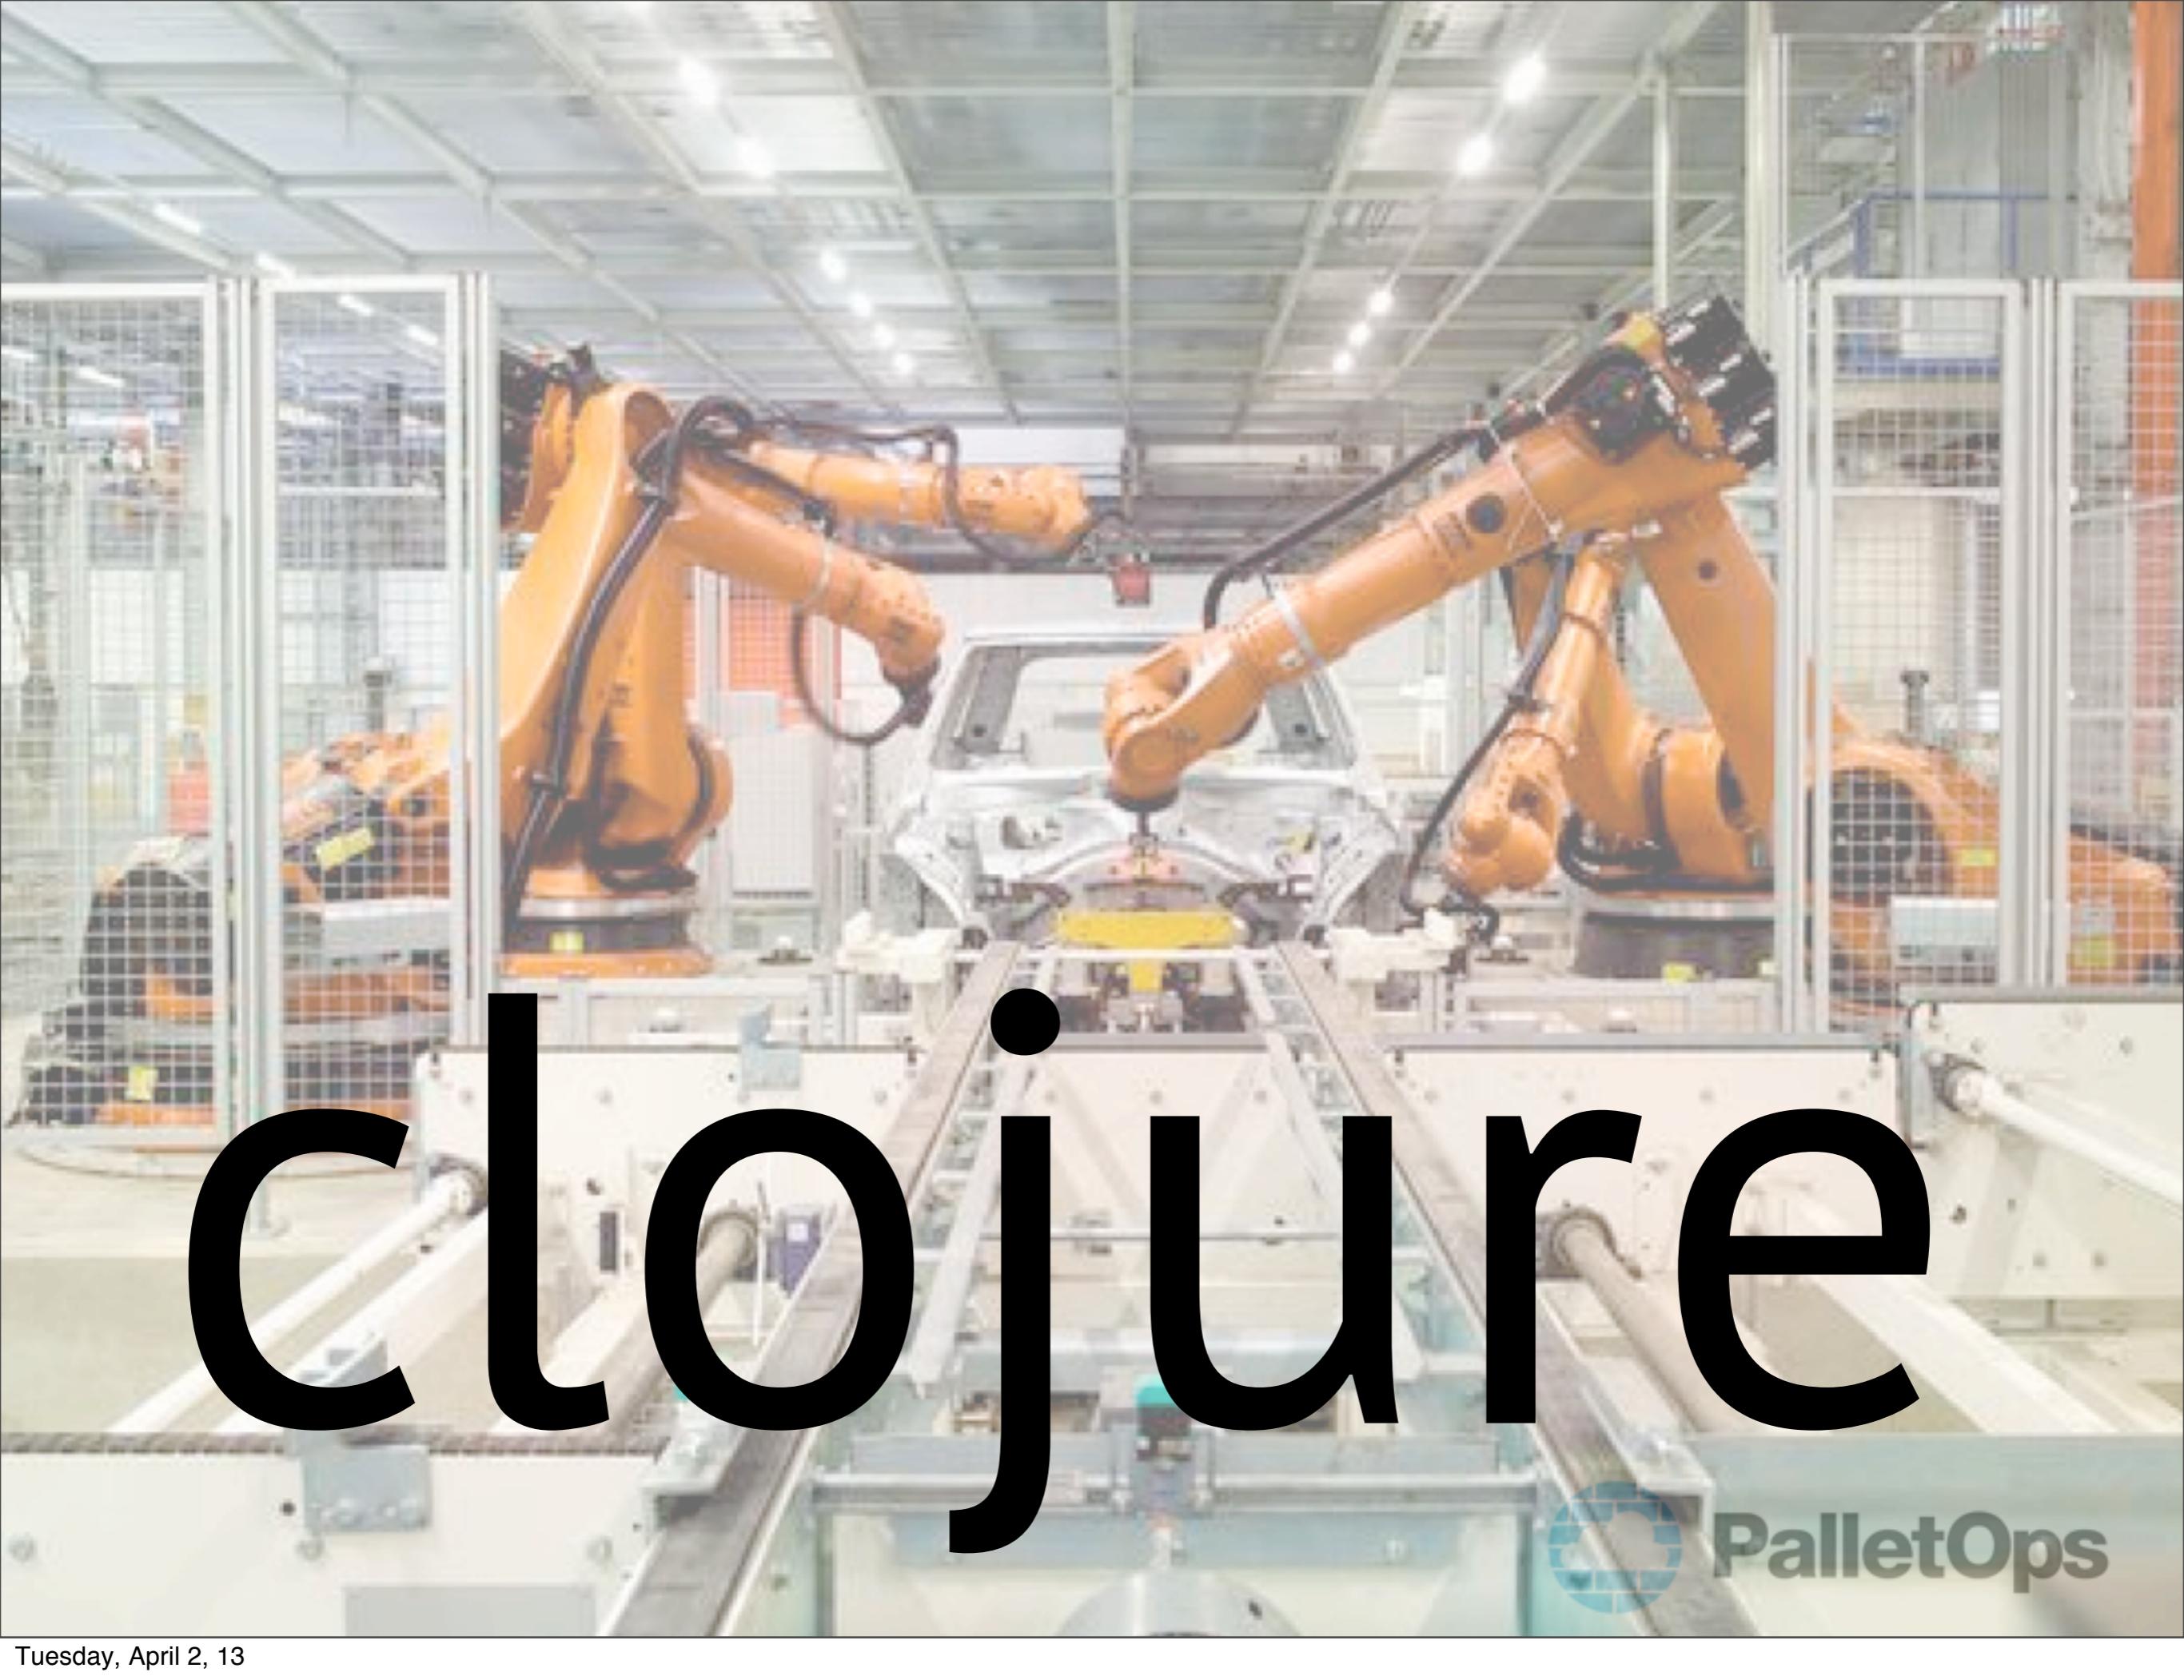




Moving Parts

(**os-families** × **os-versions** × **node-types**
× **configurations** × **cross-node configs**
× **branches** × **environments**
× **cluster shapes**
× **infrastructure providers ...**)
× **time**





clojure



PalletOps

Pallet

- An Functional Infrastructure Automation platform written in Clojure
- 3+ years of development
- Brain child of Hugo Duncan
- Used in IaaS, PaaS, SaaS and elsewhere
- +30K lines of Clojure



Design Constraints

- **Needs to work today**
- **Write once; run everywhere**
- **Native support for Cloud infrastructure**
- **Native support for Clusters**
- **A library, not a framework, not a server**
- **Build your own automation**
- **Build your own abstractions**



Pallet

- **A Library**
 - Stateless
 - Embeddable
- **Provides abstractions over**
 - Cloud Providers
 - Hardware
 - OS
 - Configuration
 - Relationships
- **A companion library of Crates (modules) a number of services**



```
(def web-server-nodes
  (node-spec {:image {:os-family :ubuntu
                      :os-version "10.04"}}
             {:hardware {:cpu-count 12
                         :min-ram (* 64 1024)}}))

(def web-servers
  (group-spec "web-server"
    :node-spec node-spec
    :phases
      {:configure (plan-fn
                    (java/install)
                    (tomcat/install)))))

(def ec2 (compute-service :ec2 ...))

(converge {web-servers 5} :compute-service ec2)
```



```
(defn web-server-node [cpus ram os-family os-version]
  (node-spec
    {:image {:os-family os-family
              :os-version os-version}}
    {:hardware {:cpu-count cpus
                :min-ram (* ram 1024)}})

(defn web-servers [cpus ram os-family os-version]
  (group-spec
    :node-spec node-spec
    :phases
      {:configure (plan-fn
                    (java/install)
                    (tomcat/install))))
```



```
(converge {(:web-servers 12 32 :centos "6.3") 1}
          :compute-service ec2)

(def platforms [[:centos "6.3"]
               [:ubuntu "10.04"]
               [:rhel "7"]])

(def webservers-to-build
  (zipmap (map (fn [[os-family os-version]]
                 (web-servers 12 32
                               os-family os-version))
               platforms)
          (repeat 1)))

(converge webservers-to-build
          :compute-service ec2)

(converge webservers-to-build
          :compute-service virtualbox)
```



Pallet

- **Serverless – your infrastructure in Git**
- **Build your abstractions on top of Pallet's abstractions**
- **Embed Pallet in your application**
- **Distribute automation as .jars**
- **Describe and build your infrastructure in a generative way**
- **Write once, run everywhere**



DSLs

- Bridging between very different worlds
- Pallet has a DSL named Stevedore for abstracting over scripting languages and tooling
 - Clojure + Stevedore + OS family/version = correct script
 - Intertwine clojure code with shell scripts in a way similar to macros do with clojure.
 - Pallet itself is built on Stevedore, and all of it is OS independent
 - Script abstractions – create and reuse scripts inside stevedore



DEMO



Multi-methods

- For when the same abstract concept has many concrete representations
- Open extensibility – very important!
- Pallet uses them for abstracting over OS family & version, at both script and plan levels, e.g.:
 - Passwordless sudo is done differently in Centos 5.3 and Centos 6+
 - Java is installed different depending on the OS family and version
 - Java is installed differently if it's Oracle JDK or OpenJDK



Data

- **Data allows to perform heavy lifting operations very simply**
- **Data is easy to test, inspect, debug, log**
- **Defer execution as much as possible**
- **Pallet internals are built around data manipulation**
 - Coupling between components is data
 - All intermediate representation is data, until right before the execution



Data

- **Pallet represents abstractions as data structures: they compose well**
 - Abstractions built as maps
 - Infrastructure can be defined dynamically and programmatically
 - Mixins with different merging semantics
 - E.g. if you merge different specs with the same function, the actions are first joined and then reordered
 - environment, node, server, group, cluster...



Compilers

- **Powerful design pattern to bridge two different worlds**
- **In a compiler**
 - Start with your code and target architecture
 - Build abstract data representation of the code
 - Live in the data world for as long as you can, transforming this data (pipeline)
 - translate to data into OPCodes for the target CPU/Architecture



plan=lift(current
system**)**

plan=lift(current
system)

desired
system =exec(plan)

DEMO



Finite State Machines

- Provide a better understanding and management of external stateful resources
- Separate state from execution
- Pallet's main execution paths are controlled by state machines
 - Separate side-effect free from side-effect all-the-things
 - Query the execution state
 - Write your own functions on existing primitives



Concurrency

- **Immutability in Clojure is priceless...**
 - ... for everything else, there are atoms and refs
- **99% of pallet is stateless**
 - functions are easy to test
 - we use threads at the target-node level
 - a session is kept in thread-local atoms
 - 100's of nodes



JVM

- **Access to well tested, well understood and high-performance libraries**
- dm;hc
- **Pallet uses 3 java libs:**
 - Multi-cloud support via jclouds (optional)
 - slf4j for logging
 - jsch – ssh support



JVM

- **.jar**
 - Embed Pallet in your code
 - Distribute pallet automation as jars
 - Build your domain on Pallet
 - Every developer can build their own infrastructure
 - Projects can distribute their infrastructure builder
 - Versioning, signing, etc...



$b \approx f(a) ?$

Where are we now?

- Functional and programmatic infrastructure automation
- Works on most cloud providers and target OSs (as long as they're *nix)
- Build complex and flexible clusters
- Fast development paths
- Easy to build your domain abstractions on infrastructure



What are we doing with it?

- **Pallet Big Data – Automate everything Big Data**
 - Amazon EMR everywhere!
 - Build Clusters, keep them happy, run workloads
- **bricklayer – like git, but for builds**
 - Decouple build from server
 - Build evolves with code
- **lein pallet up – download project from git, run it everywhere**





PalletOps

Infrastructure Automation

~

Clojure Development

~

<http://palletops.com>