

The life and times of

PuppetDB

Clojure/West 2013

deepak
giridharagopal

deepak@puppetlabs.com
@grim_radical [github twitter freenode]

Let's talk about...

~~Sorry~~ NO
CHANGE

Immutability
is great!

Immutability allows
for invariants, which
help you reason about
correctness

**Immutability prevents
spooky action at a
distance**

**Immutability fosters
modular, composable
abstractions**

*(not a tough sell
to Clojurists*)*

**clojurians?*

That's great for
development, but how
about operations?

Immutability for infrastructure?

*Because operations is in the
same boat as development*

**Everyone who's got
their app running on
a fleet of servers has
experienced **spooky**
action at a distance**

Known, good state is
critical for reliable
upgrades

A lack of predictability
in your systems
**ruins automation and
abstraction**

The problem is that:

**Systems are inherently
mutable!**

But ideally:

**Systems should behave as
though they weren't!**

façade of immutability

Computer systems are in many ways open systems, providing the keys to the vault if one is so inclined to grab them. But in order to foster an air of immutability in our own systems, it's of utmost importance to create a façade of immutability. Immutability requires that we layer over and abstract the parts of our system that provide unrestrained mutability.



Describe how you'd
like your systems to
look, and **Puppet**
does all the hard
work for you!

```
file { “/etc/issue”:
  content => “Got an issue? Here’s a tissue!”,  
}
```

```
file { “/etc/motd”:
  content => template(“Welcome to $hostname!”),  
}
```

```
file { "/etc/sudoers":  
    owner  => root,  
    group  => root,  
    mode    => 440,  
    source  => "puppet:///modules/sudo/sudoers"  
}
```

```
package { 'ntp':
  ensure => installed,
}

service { 'ntpd':
  ensure     => running,
  enable     => true,
  subscribe => File['/etc/ntp.conf'],
}

file { '/etc/ntp.conf':
  ensure  => file,
  require => Package['ntp'],
  source   => "puppet:///modules/ntp/ntp.conf",
}
```

```
class ntp {  
  
    package { 'ntp':  
        ensure => installed,  
    }  
  
    service { 'ntpd':  
        ensure      => running,  
        enable      => true,  
        subscribe  => File['/etc/ntp.conf'],  
    }  
  
    file { '/etc/ntp.conf':  
        ensure  => file,  
        require => Package['ntp'],  
        source   => "puppet:///modules/ntp/ntp.conf",  
    }  
}
```

```
node "webserver.mydomain.com" {  
    includentp  
}
```

```
node "appserver.mydomain.com" {  
    includentp  
}
```

```
node "database.mydomain.com" {  
    includentp  
}
```

```
class ssh {  
  
    @@sshkey { $hostname:  
        type => dsa,  
        key => $sshdssakey  
    }  
  
    Sshkey <<| |>>  
}  
}
```

File “/tmp/foo/bar”

User “deepak”

Dir “/tmp/foo”

Dir “/tmp”

Dir “/tmp”

User “deepak”

Dir “/tmp/foo”

File “/tmp/foo/bar”



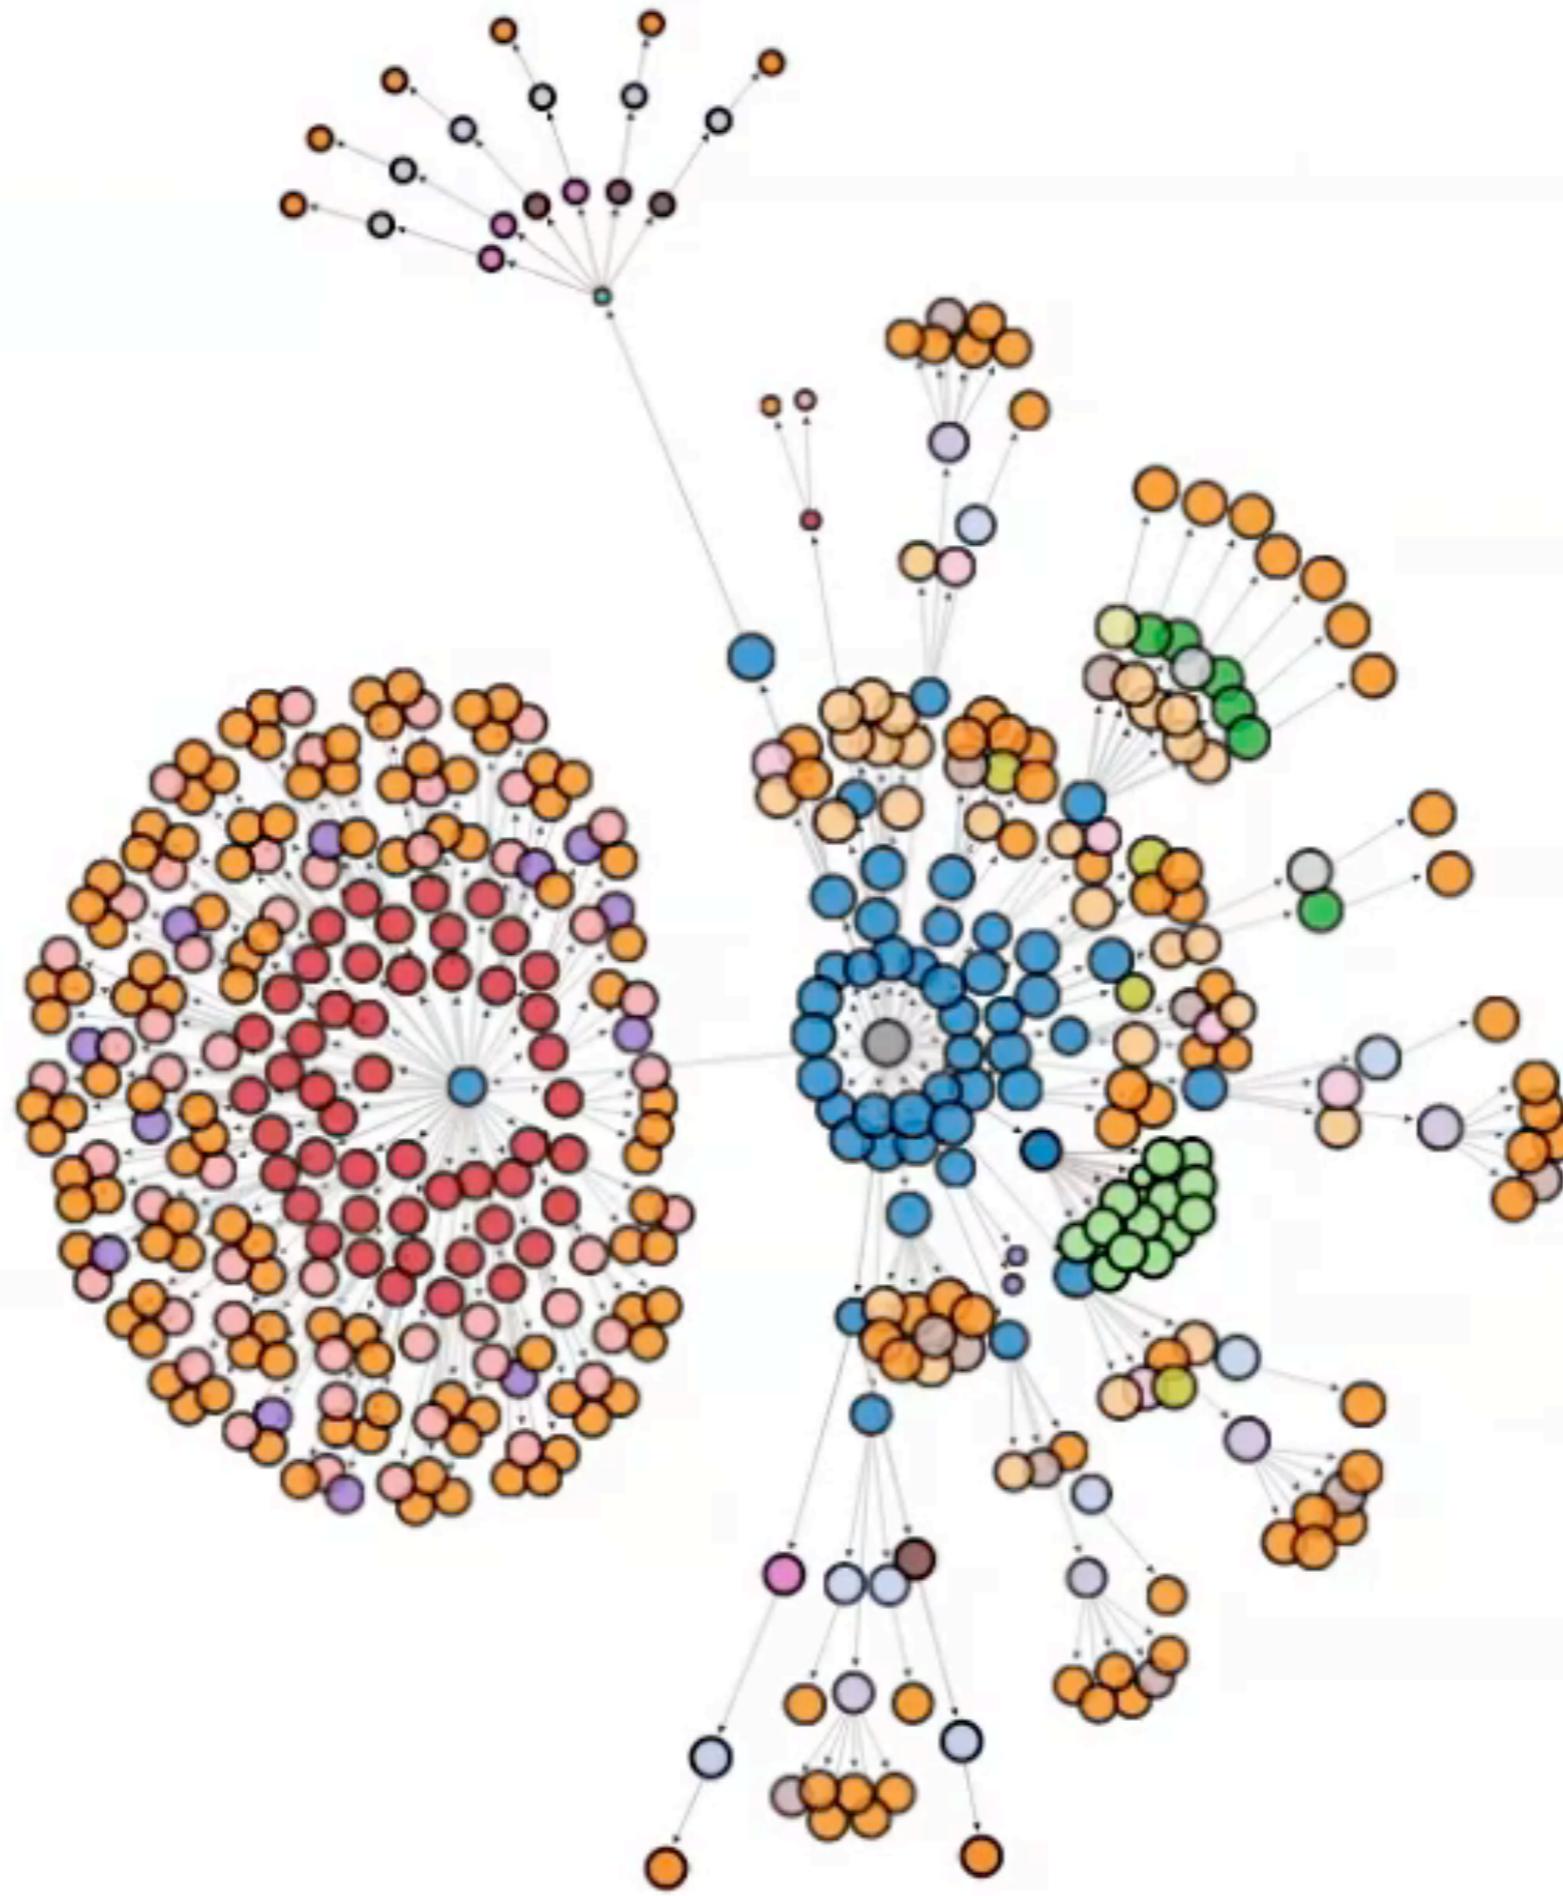
Dir “/tmp”

User “deepak”

Dir “/tmp/foo”

File “/tmp/foo/bar”





**Idempotent, and
only does what's
necessary**

**Compensates for the
inherent mutability
of systems**

**Combats spooky
action at a distance
with automatic
repair**

**Brings predictability
to your systems**

A foundation of predictability and reliability lets you perform **higher-level operations** on your infrastructure

vOLUME

**every resource
every parameter
every relationship
every fact**

**for every node
updated all the time**

**Users leverage this
data to do higher-
order things with
their infrastructure**

key distribution
monitoring
clustered services
master/slave replication
load balancers
shared filesystems
firewall rules

...

Infrastructure as code

Infrastructure as **data**

User demand:

Store as much data as we can!
Much better queryability!

Oh yeah, but:

Don't slow down the system!
Don't compromise reliability!

We can rebuild it, we
have the technology!

Speed is important

Parsing, validating, and manipulating incoming data is computationally expensive

Speed is important

The slower central storage is, the less agile sysadmins can be. That can cost money and uptime!

Reliability is important

If it's a critical part of managing infrastructure, it's got to be solid!

Deployment is important

*Should be easy to package,
install, configure, use, monitor,
and upgrade.*

**Wait, isn't Puppet
written in Ruby?**

Lessons learned from writing the rest of our software in Ruby

*It's...not speedy. Object creation,
method calls, garbage collection,
etc. “Magical” APIs amplify
the problem.*

Lessons learned from writing the rest of our software in Ruby

*Can only use one core!
Workarounds compromise
performance or simplicity*

Lessons learned from writing the rest of our software in Ruby

*Mutable state all over
the damn place!*

Lessons learned from writing the rest of our software in Ruby

*Many struggles with the
runtime. :(*



-- Jeff Gagliardi





1. It's fast

2. JVM libraries & Tools

3. State & Parallelism

PuppetDB

PuppetDB

Definitely Better!

**Fast, safe storage
of catalogs, facts,
and events**

*like, *way* faster!*

HTTP APIs

for resource, fact,
node, report retrieval

*plenty of data, just
a “curl” away!*

Storage & Querying

Storage & Querying

CORS

Command Query Responsibility Separation

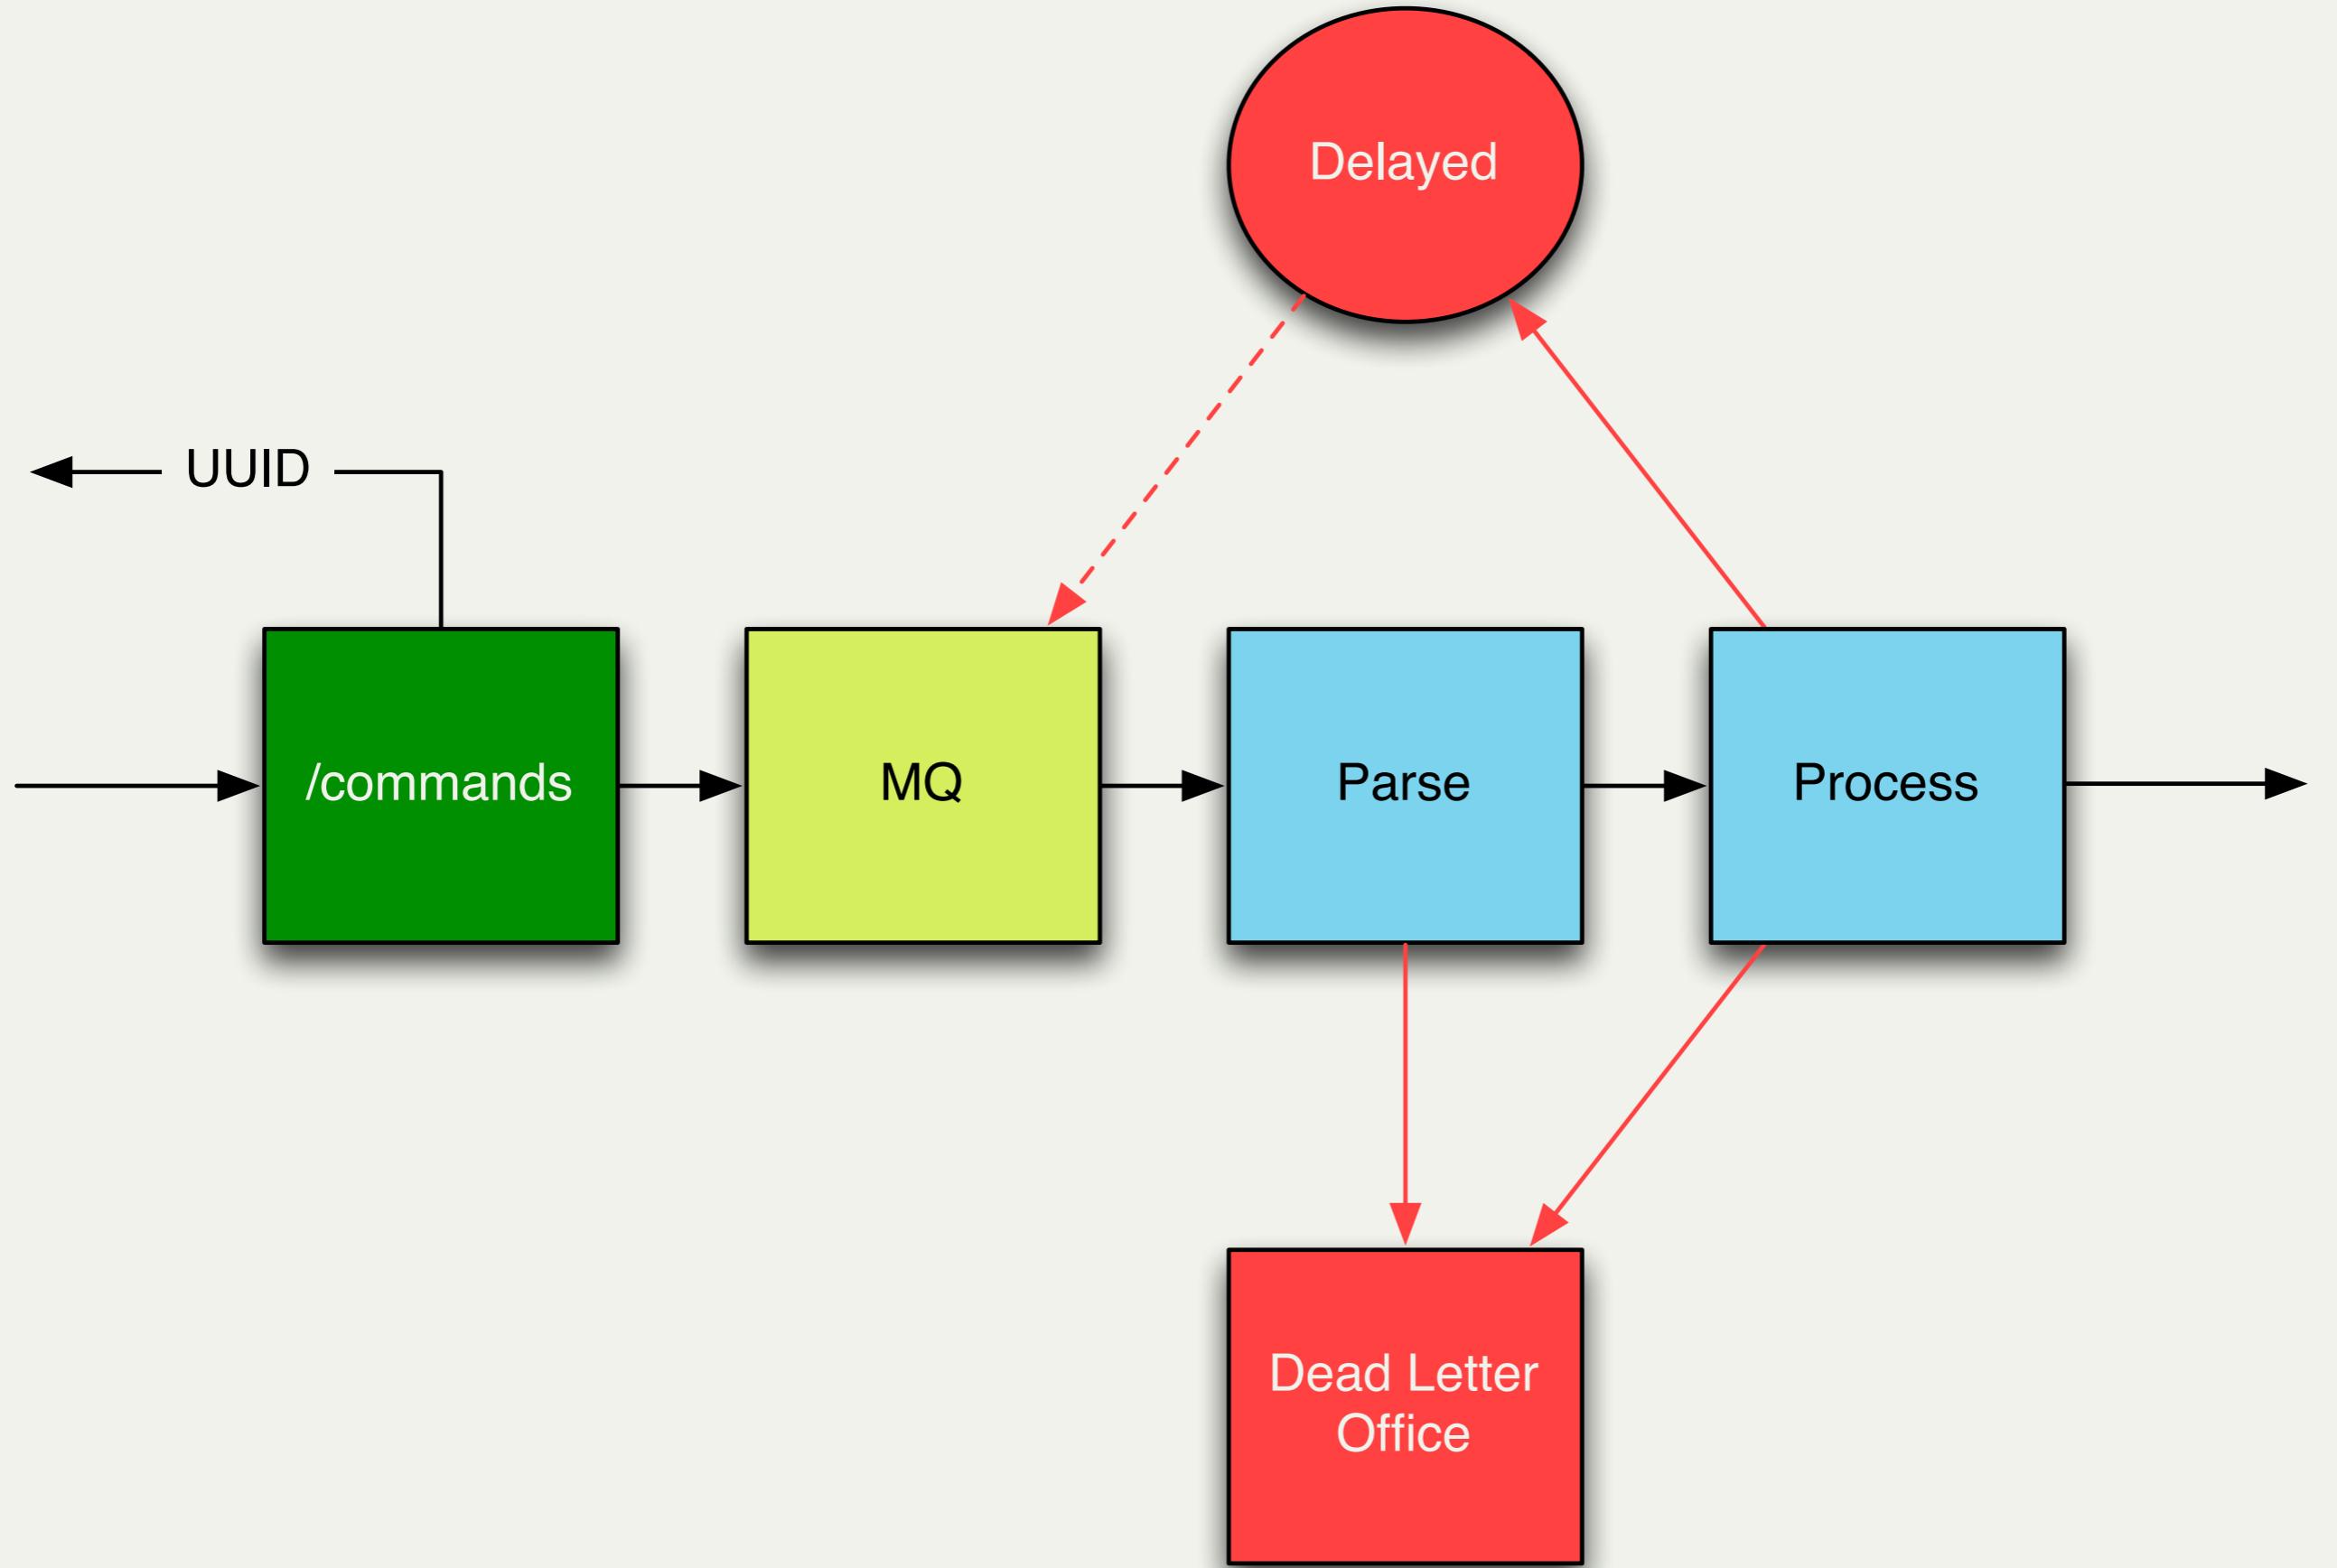
*use a different model to update
information than the model you
use to read information*

Writes

cqrs write pipeline

*async, parallel, MQ-based, with
automatic retry*

```
{  
  :command "replace catalog"  
  :version 2  
  :payload {...}  
}
```



```
(defmulti process-command!
  (fn [{:keys [command version] :or {version 1}} _]
    [command version]))  
  
(defmethod process-command! ["replace catalog" 1]
  [command options]
  (replace-catalog* command options))
```

```
(defmulti process-command!
  (fn [{:keys [command version] :or {version 1}} _]
    [command version]))  
  
(defmethod process-command! ["replace catalog" 2]
  [command options]
  (replace-catalog* command options))  
  
(defmethod process-command! ["replace catalog" 1]
  [command options]
  (-> command
      (update-in [:payload] catalog-v1->v2)
      (replace-catalog* options)))
```

Command
processors must be
retry-aware

*expect failure, because
it *will* happen.*

(demo)

Reads

Queries

are expressed in their
own “language”

*domain specific, AST-based
query language*

```
[ "and",
  [ "=", "type", "User" ],
  [ "=", "title", "nick" ] ]
```

```
["and",
 ["=", ["fact", "operatingsystem"], "Debian"],
 ["<", ["fact", "uptime_seconds"], 1000]]
```

```
["and",
 ["=", "name", "ipaddress"],
 ["in", "certname",
  ["extract", "certname", ["select-resources",
    ["and",
     ["=", "type", "Class"],
     ["=", "title", "Apache"]]]]]
```

```
["or",
  ["=", "certname", "foo.com"],
  ["=", "certname", "bar.com"],
  ["=", "certname", "baz.com"]]
```

We use core.match
to walk the tree,
compiling it to SQL

AST-based API lets
users **write their own**
languages

*ah, you've got to love
open source!*

```
(Package[httpd] and country=fr)  
or country=us
```

```
Package["mysql-server"]  
and architecture=amd64
```

Erik Dalén, Spotify

<https://github.com/dalen/puppet-puppetdbquery>

**AST-based API lets
us more safely
manipulate queries**

```
(def query-app
  (app
    [&]
    {:get (fn [{:keys [params] :as request}]
            (perform-query (params "query"))))}))
```

```
(def resources-app
  (app
    []
    query-app

    [type title &]
    (comp query-app
          (partial http-q/restrict-resource-query-to-type type)
          (partial http-q/restrict-resource-query-to-title title)))

    [type &]
    (comp query-app
          (partial http-q/restrict-resource-query-to-type type))))
```

```
(defn restrict-resource-query-to-type
  [type query]
  (conj [ "and"
          [ "=" "type" type] ]
        query) )
```

Storage

Relational Database, embedded or PostgreSQL

*because they're actually pretty
fantastic at ad-hoc queries,
aggregation, windowing, etc.
while maintaining safety*

Fault-tolerance

by @jrecursive



Relational Database, embedded or PostgreSQL

*we use arrays, recursive queries,
indexing inside complex
structures*

Relational Database, embedded or PostgreSQL

*schema is oriented towards the
types of queries we encounter*

Deployment

PuppetDB Server

DLO

DB

Workers

HTTP

MQ

PuppetDB Server

Workers

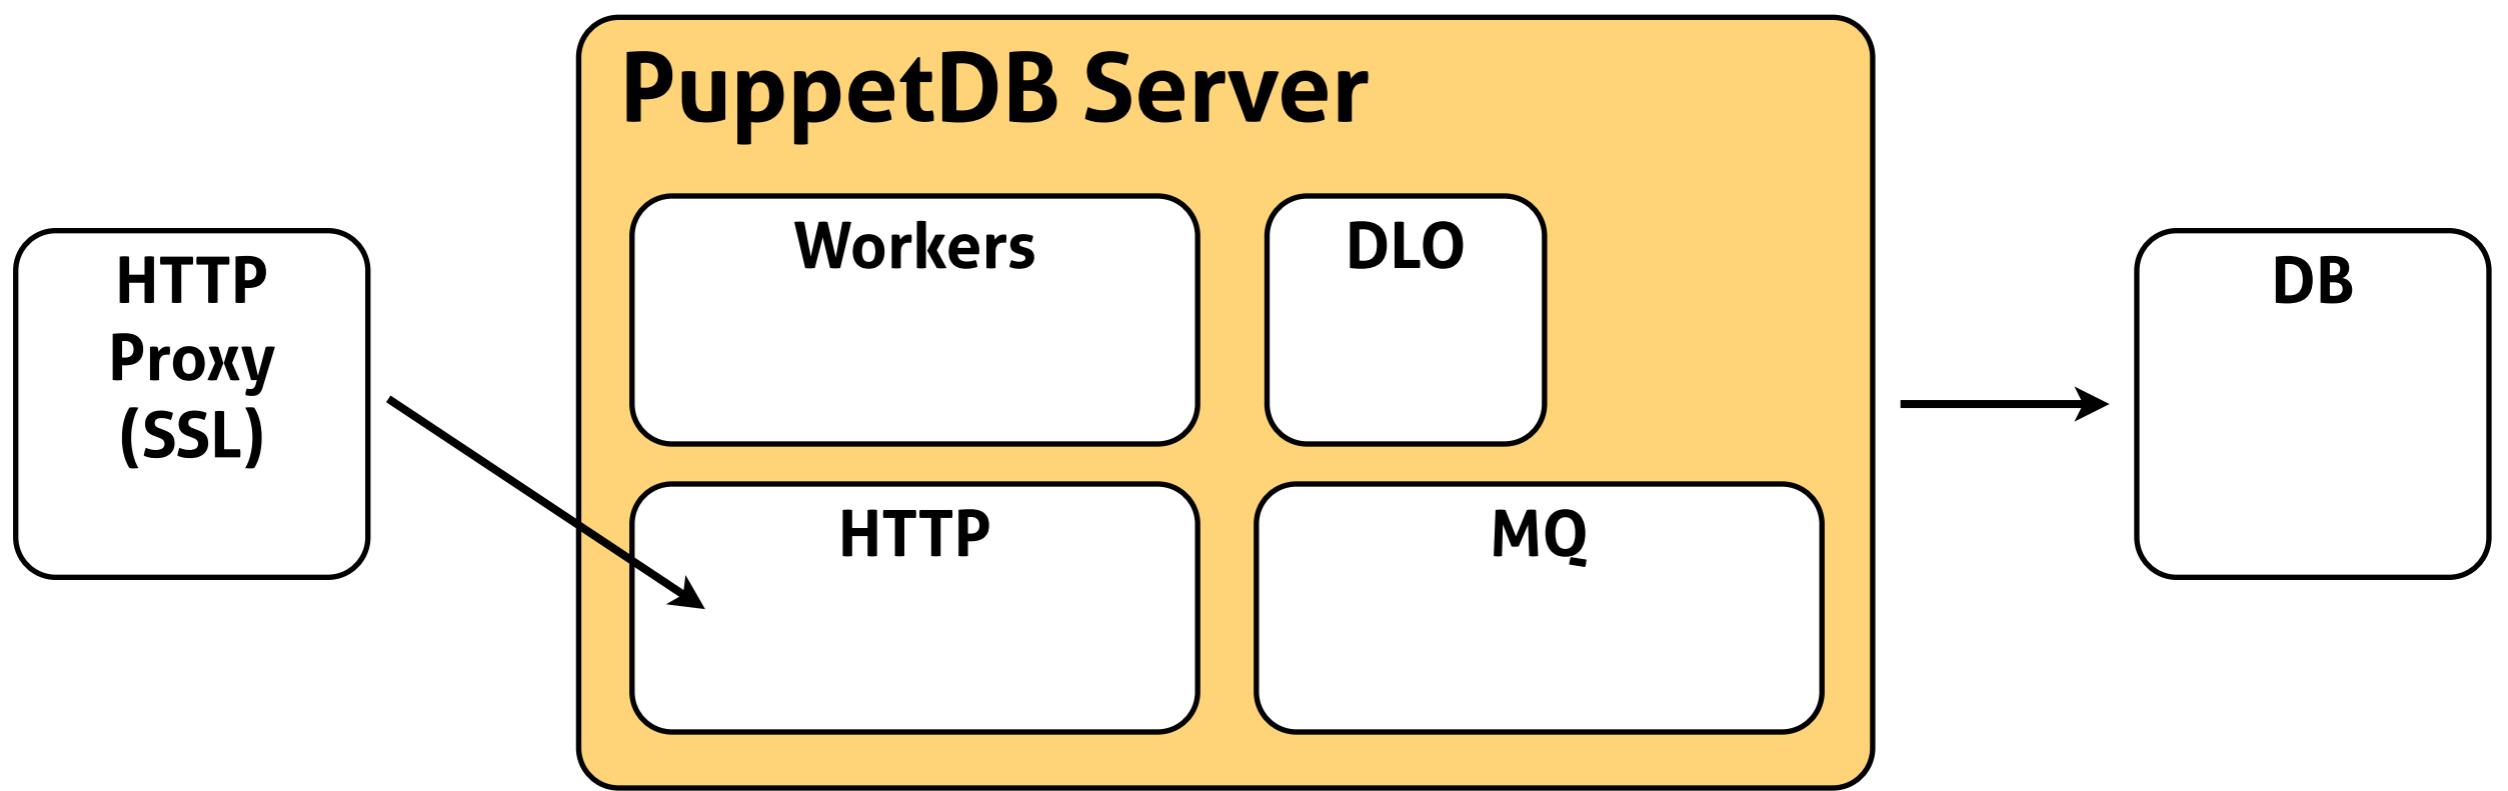
DLO

HTTP

MQ

DB





Codez

```
(defn process-commands!
  "Connect to an MQ and continually, sequentially process commands.

  If the MQ consumption timeout is reached without any new data, the
  function will terminate."
  [connection endpoint discarded-dir options-map]
  (let [producer    (mq-conn/producer connection)
        publish     (partial mq-producer/publish producer endpoint)
        on-message  (produce-message-handler publish discarded-dir options-map)
        mq-error    (promise)
        consumer    (mq-conn/consumer connection
                                         {:endpoint endpoint
                                          :on-message on-message
                                          :transacted true
                                          :on-failure #(deliver mq-error (:exception %))}))]
    (mq-cons/start consumer)

    ;; Block until an exception is thrown by the consumer thread
    (try
      (deref mq-error)
      (throw @mq-error)
      (finally
        (mq-cons/close consumer)))))
```

```
(defmacro with-error-delivery
  "Executes body, and delivers an exception to the provided promise if one is
  thrown."
  [error & body]
  `(~@body
    (try
      (catch Exception e#
        (deliver ~error e#))))
```

```
(defn filter-mbean
  "Converts an mbean to a map. For attributes that can't be converted to JSON,
  return a string representation of the value."
  [mbean]
  {:post [(map? %)]}
  (into {} (for [[k v] mbean]
              (cond
                ;; Nested structures should themselves be filtered
                (map? v)
                [k (filter-mbean v)]

                ;; Cheshire can serialize to JSON anything that
                ;; implements the JSONable protocol
                (satisfies? JSONable v)
                [k v]

                :else
                [k (str v)])))))
```

```
(defmacro with-test-broker
  "Constructs and starts an embedded MQ, and evaluates `body` inside a
  `with-open` expression that takes care of connection cleanup and MQ
  tear-down."
```

`name` - The name to use for the embedded MQ

`conn-var` - Inside of `body`, the variable named `conn-var` contains an active connection to the embedded broker.

Example:

```
(with-test-broker \"my-broker\" the-connection
  ;; Do something with the connection
  (prn the-connection))

[ name conn-var & body]
`(let [dir#
       broker-name#
       conn-str#
       ^BrokerService broker#]
    (fs/absolute-path (fs/temp-dir))
    ~name
    (str "vm://" ~name)
    (mq/build-embedded-broker broker-name# dir#) ]

  (.setUseJmx broker# false)
  (.setPersistent broker# false)
  (mq/start-broker! broker#)

  (try
    (with-open [<conn-var (mq/connect! conn-str#)]
      ~@body)
    (finally
      (mq/stop-broker! broker#)
      (fs/delete-dir dir#))))
```

Results

Thousands of production deployments

Small shops with a dozen hosts,
large shops with thousands of
hosts, intercontinental
deployments...



Dean Wilson @unixdaemon

Did some basic performance testing of #puppetdb and reduced our compile times from 180ish second to mid 20s. #massivewin

7h

[Expand](#)



Greg Mason @greg_mason

25 Sep

according to #puppetdb, I'm currently managing 71,777 resources with Puppet.

[Expand](#)



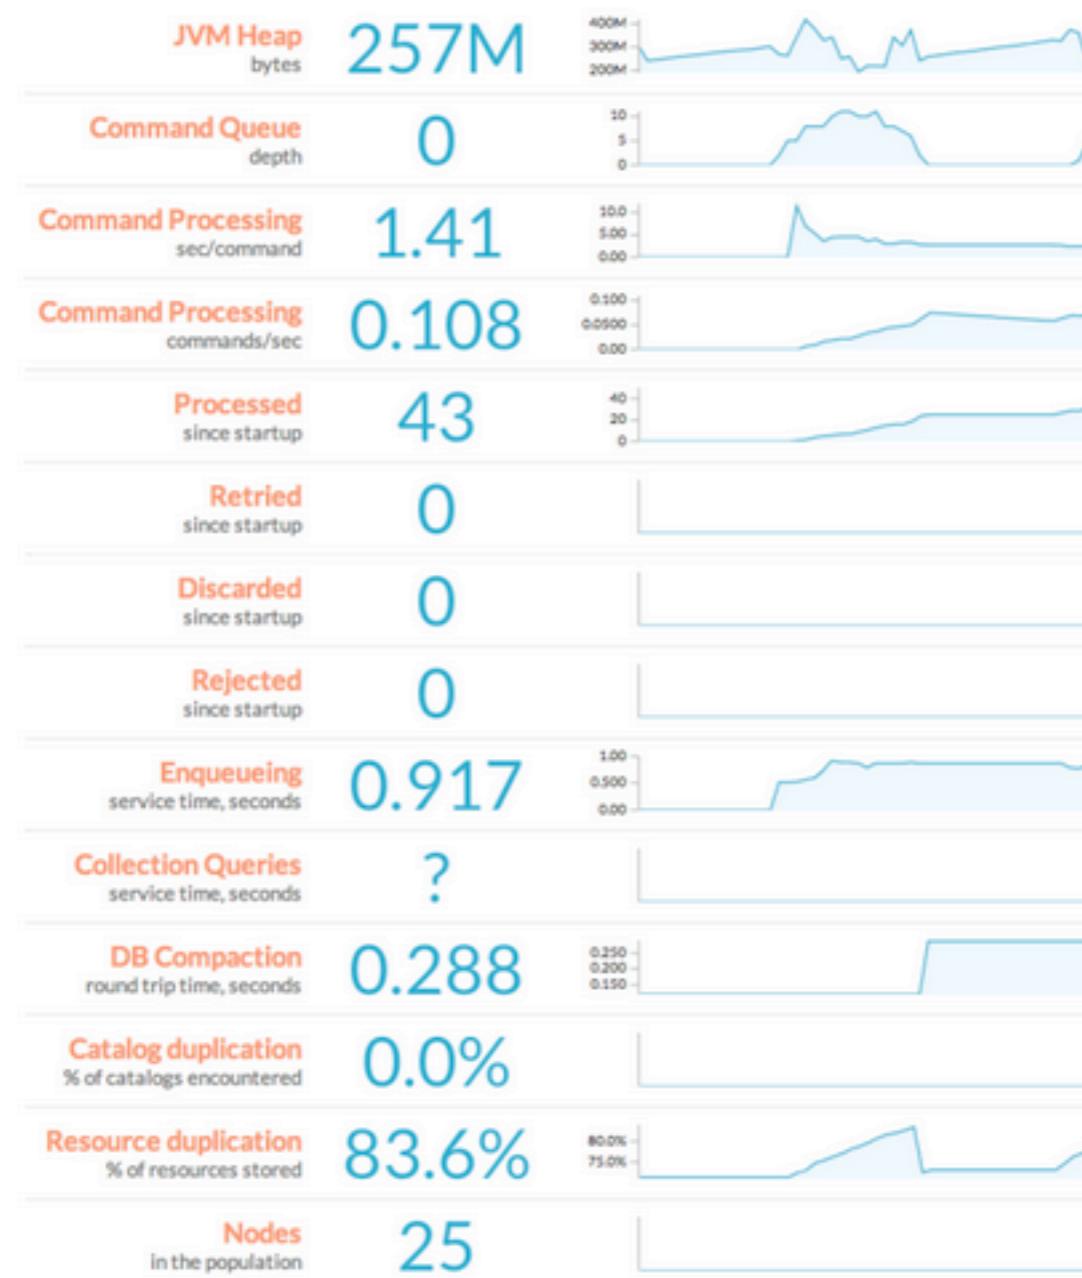
ripienaar @ripienaar

13 Jun

puppetdb + 25 masterless nodes with embedded database
twitpic.com/9vz7y2

Retweeted by D Giridharagopal

Hide photo Reply Retweeted Favorite



Matt Nicholson @sjoebboo

@glarizza @grim_radical two pic.twitter.com/mvJuOijF

Hide photo



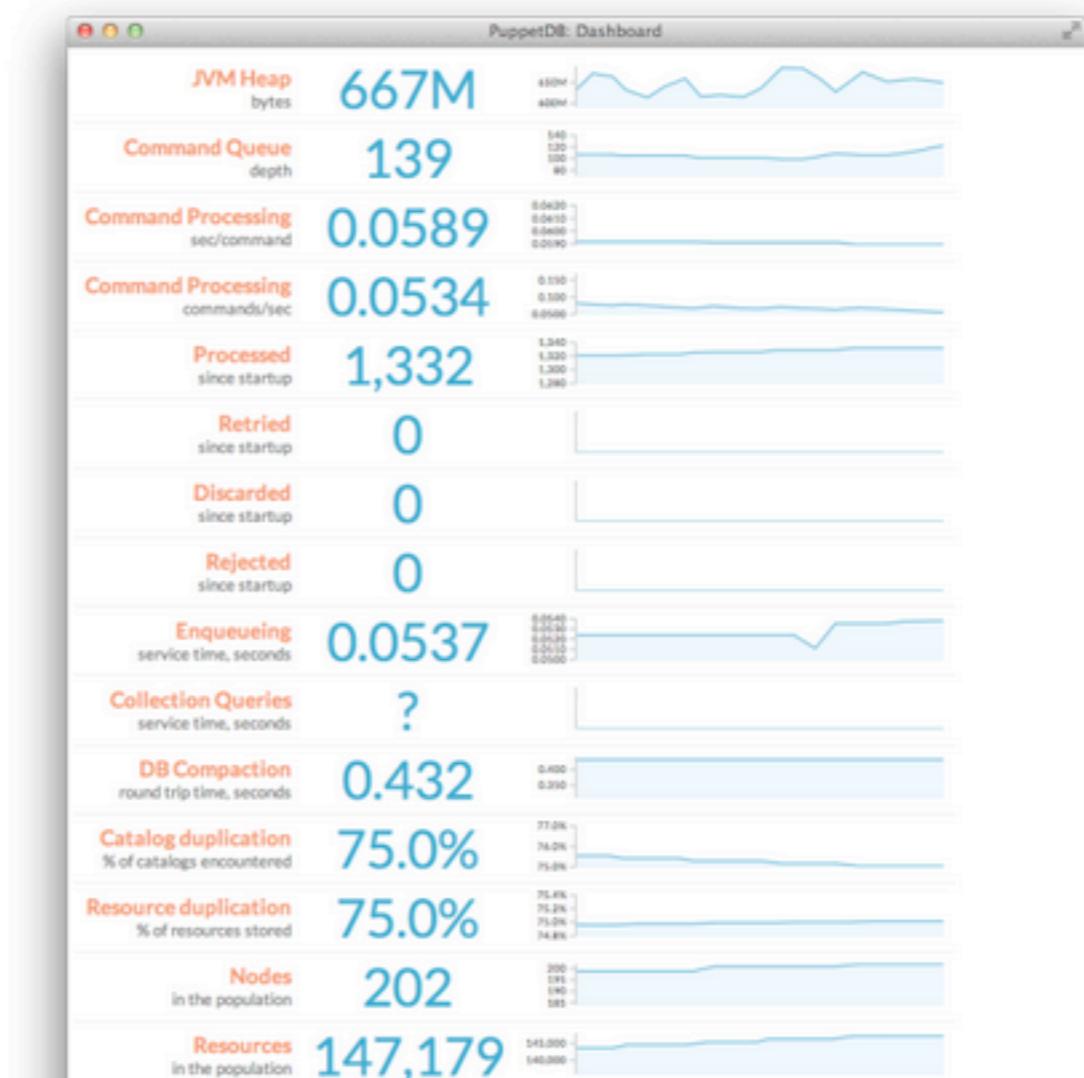
Lars Fronius @LarsFronius

23 May

This #puppetdb dashboard thingy is actually so awesome. Oh, and puppetdb is such a big step in performance! #puppetize
twitpic.com/9oi290

Retweeted by D Giridharagopal

Hide photo Reply Retweeted Favorite



20 Jul

Thousands of deployments,
Hundreds of threads per install,
Zero deadlocks,
Zero bugs involving clojure state

*companion Ruby code has
~10x the defect rate*

**Users dig the uberjar,
makes deployment
straightforward and
less error-prone**

**Nobody cares that it's
Clojure. Users just
want stuff to work.**

**Monitoring and
instrumentation is a
big deal. Users want
easy ways to
consume metrics.**

**If you build it, they
will come.**

Open source

[http://github.com/
puppetlabs/puppetdb](http://github.com/puppetlabs/puppetdb)

deepak
giridharagopal

deepak@puppetlabs.com
@grim_radical [github twitter freenode]

We're hiring!